

Batch loading data

Daniel Anderson

Week 4, Class 1

Agenda

- Discuss the midterm
 - Canvas quiz (10 points; please don't stress)
 - Take home (40 points)
- Review Lab 1
- `map_dfr` and batch-loading data
- Introduce list columns
 - This will mostly be an intro to Wednesday's lecture

Learning objectives

- Understand when `map_dfr` can and should be applied
- Better understand file paths, and how `{fs}` can help
- Be able to batch load data of a specific type within a mixed-type directory
- Use filenames to pull data

Midterm

Questions?

Let's look at the take-home portion

Review Lab 1

map_dfr

- If each iteration returns a data frame, you can use `map_dfr` to automatically bind all the data frames together.

Example

- Create a function that simulates data (please copy the code and follow along)

```
set.seed(8675309)
simulate <- function(n, mean = 0, sd = 1) {
  tibble(sample_id = seq_len(n),
          sample = rnorm(n, mean, sd))
}
simulate(10)
```

```
## # A tibble: 10 x 2
##   sample_id      sample
##   <int>         <dbl>
## 1         1 -0.9965824
## 2         2  0.7218241
## 3         3 -0.6172088
## 4         4  2.029392
## 5         5  1.065416
## 6         6  0.9872197
## 7         7  0.02745393
## 8         8  0.6728723
## 9         9  0.5720665
## 10        10  0.9036777
```

Two more quick examples

```
simulate(3, 100, 10)
```

```
## # A tibble: 3 x 2
##   sample_id    sample
##   <int>      <dbl>
## 1         1  84.50448
## 2         2 110.2264
## 3         3 101.5008
```

```
simulate(5, -10, 1.5)
```

```
## # A tibble: 5 x 2
##   sample_id    sample
##   <int>      <dbl>
## 1         1 -10.98995
## 2         2 -11.49188
## 3         3  -7.041312
## 4         4 -10.66270
## 5         5 -11.35096
```


Simulation

- Assume we want to vary the sample size from 10 to 150 by increments of 5
- `mean` stays constant at 100, `sd` is constant at 10

Try with `purrr::map`

02:00

```
library(tidyverse)
sims <- map(seq(10, 150, 5), simulate, 100, 10)
```

```
sims[1]
```

```
## [[1]]
## # A tibble: 10 x 2
##   sample_id sample
##   <int>     <dbl>
## 1         1 103.7618
## 2         2 111.5353
## 3         3 115.7490
## 4         4 105.8853
## 5         5  93.84955
## 6         6  97.71089
## 7         7 100.6392
## 8         8  96.86526
## 9         9  97.51501
## 10        10  98.46205
```

```
sims[2]
```

```
## [[1]]
## # A tibble: 15 x 2
##   sample_id sample
##   <int>     <dbl>
## 1         1  93.64743
## 2         2  99.96206
## 3         3 100.4562
## 4         4 106.8407
## 5         5  97.47957
## 6         6  98.48961
## 7         7  91.25069
## 8         8  80.23099
## 9         9 102.3766
## 10        10 100.3609
## 11        11 101.3490
## 12        12 101.1758
## 13        13  91.74411
## 14        14  78.64764
## 15        15 102.1421
```

Swap for `map_dfr`

Try it – what happens?

```
sims_df <- map_dfr(seq(10, 150, 5), simulate, 100, 10)
sims_df
```

```
## # A tibble: 2,320 x 2
##   sample_id    sample
##   <int>      <dbl>
## 1         1  85.64361
## 2         2 103.6789
## 3         3  94.71782
## 4         4 103.1350
## 5         5  99.78701
## 6         6 105.3462
## 7         7 100.0653
## 8         8  94.28314
## 9         9 108.8872
## 10        10 106.0850
## # ... with 2,310 more rows
```

01:00

Notice a problem here

```
sims_df[1:15, ]
```

```
## # A tibble: 15 x 2
##   sample_id    sample
##   <int>      <dbl>
## 1         1  85.64361
## 2         2 103.6789
## 3         3  94.71782
## 4         4 103.1350
## 5         5  99.78701
## 6         6 105.3462
## 7         7 100.0653
## 8         8  94.28314
## 9         9 108.8872
## 10        10 106.0850
## 11         1  89.49968
## 12         2  86.99898
## 13         3  85.38054
## 14         4  99.10690
## 15         5 105.0088
```

.id argument

```
sims_df2 <- map_dfr(seq(10, 150, 5), simulate, 100, 10,  
                   .id = "iteration")  
sims_df2[1:14, ]
```

```
## # A tibble: 14 x 3  
##   iteration sample_id      sample  
##   <chr>      <int>      <dbl>  
## 1 1          1 112.1250  
## 2 1          2  88.07056  
## 3 1          3 108.3908  
## 4 1          4 100.8193  
## 5 1          5 102.1545  
## 6 1          6 113.5398  
## 7 1          7 101.4171  
## 8 1          8  99.33668  
## 9 1          9 100.2855  
## 10 1         10  90.22043  
## 11 2          1  91.08882  
## 12 2          2 107.3664  
## 13 2          3 101.1745  
## 14 2          4  96.82053
```

.id: Either a string or NULL. If a string, the output will contain a variable with that name, storing either the name (if `.x` is named) or the index (if `.x` is unnamed) of the input. If NULL, the default, no variable will be created.

– `{purrr}` documentation

setNames

```
sample_size <- seq(10, 150, 5)
sample_size
```

```
## [1] 10 15 20 25 30 35 40 45 50 55 60 65 70 75 80 85 90
## [23] 120 125 130 135 140 145 150
```

```
sample_size <- setNames(sample_size,
                        english::english(seq(10, 150, 5)))
sample_size[1:15]
```

```
##          ten          fifteen          twenty  twenty-five          thirty  thirty
##          10           15           20           25           30
##  forty-five        fifty  fifty-five        sixty  sixty-five        se
##          45           50           55           60           65
##      eighty
##          80
```

Try again

```
sims_df3 <- map_dfr(sample_size, simulate, 100, 10,  
                    .id = "n")  
sims_df3[1:14, ]
```

```
## # A tibble: 14 x 3  
##       n      sample_id      sample  
##   <chr>      <int>      <dbl>  
## 1 ten           1  98.94914  
## 2 ten           2 101.6824  
## 3 ten           3  88.16447  
## 4 ten           4  90.13604  
## 5 ten           5  85.53591  
## 6 ten           6  90.69977  
## 7 ten           7 105.8858  
## 8 ten           8  89.12978  
## 9 ten           9 114.4982  
## 10 ten          10 111.6440  
## 11 fifteen      1 103.2732  
## 12 fifteen      2 106.8949  
## 13 fifteen      3  88.83591  
## 14 fifteen      4 105.5402
```


Another quick example

broom::tidy

- The {broom} package helps us extract model output in a tidy format

```
lm(tvhours ~ age, gss_cat) %>%  
  broom::tidy()
```

```
## # A tibble: 2 x 5  
##   term          estimate std.error statistic    p.value  
##   <chr>          <dbl>    <dbl>    <dbl>    <dbl>  
## 1 (Intercept)  1.992391  0.06980966  28.54033 4.672161e-173  
## 2 age          0.02095679 0.001387361  15.10551 4.689310e- 51
```

Fit separate models by year

Again – probs not best statistically

```
split(gss_cat, gss_cat$year) %>%  
  map_dfr(~lm(tvhours ~ age, .x) %>%  
    broom::tidy())
```

```
## # A tibble: 16 x 5  
##   term          estimate std.error statistic    p.value  
##   <chr>          <dbl>      <dbl>      <dbl>      <dbl>  
## 1 (Intercept)  2.080163  0.1709061  12.17138  7.995632e-33  
## 2 age          0.01948584 0.003485199  5.591027  2.599011e- 8  
## 3 (Intercept)  2.078999  0.2176829   9.550583  1.191266e-20  
## 4 age          0.01963575 0.004400292  4.462375  9.137366e- 6  
## 5 (Intercept)  1.767990  0.2464509   7.173804  1.531756e-12  
## 6 age          0.02386070 0.005031548  4.742218  2.459650e- 6  
## 7 (Intercept)  2.096054  0.1496431  14.00702  1.419772e-42  
## 8 age          0.01781388 0.002977289  5.983256  2.589482e- 9  
## 9 (Intercept)  1.855278  0.2156381   8.603668  2.167351e-17  
## 10 age         0.02390720 0.004314567  5.541043  3.628675e- 8  
## 11 (Intercept)  2.068914  0.2096397   9.868903  2.896085e-22  
## 12 age         0.01989505 0.004086638  4.868317  1.251234e- 6  
## 13 (Intercept)  1.878070  0.2258400   8.315932  2.280108e-16  
## 14 age         0.02547794 0.004449295  5.726287  1.274840e- 8  
## 15 (Intercept)  1.980095  0.1877544  10.54620  3.238043e-25
```

.id

In cases like the preceding, `.id` becomes invaluable

```
split(gss_cat, gss_cat$year) %>%  
  map_dfr(~lm(tvhours ~ age, .x) %>%  
    broom::tidy(),  
    .id = "year")
```

```
## # A tibble: 16 x 6  
##   year term      estimate std.error statistic    p.value  
##   <chr> <chr>      <dbl>      <dbl>      <dbl>      <dbl>  
## 1 2000 (Intercept) 2.080163  0.1709061  12.17138 7.995632e-33  
## 2 2000 age      0.01948584 0.003485199  5.591027 2.599011e- 8  
## 3 2002 (Intercept) 2.078999  0.2176829   9.550583 1.191266e-20  
## 4 2002 age      0.01963575 0.004400292  4.462375 9.137366e- 6  
## 5 2004 (Intercept) 1.767990  0.2464509   7.173804 1.531756e-12  
## 6 2004 age      0.02386070 0.005031548  4.742218 2.459650e- 6  
## 7 2006 (Intercept) 2.096054  0.1496431  14.00702 1.419772e-42  
## 8 2006 age      0.01781388 0.002977289  5.983256 2.589482e- 9  
## 9 2008 (Intercept) 1.855278  0.2156381   8.603668 2.167351e-17  
## 10 2008 age      0.02390720 0.004314567  5.541043 3.628675e- 8  
## 11 2010 (Intercept) 2.068914  0.2096397   9.868903 2.896085e-22  
## 12 2010 age      0.01989505 0.004086638  4.868317 1.251234e- 6  
## 13 2012 (Intercept) 1.878070  0.2258400   8.315932 2.280108e-16  
## 14 2012 age      0.02547794 0.004449295  5.726287 1.274840e- 8
```

Batch— loading data

Please follow along

{fs}

- note – there are base equivalents. `{fs}` is just a bit better across platforms and has better defaults.

Could we apply `map_dfr` here?

```
# install.packages("fs")
library(fs)
dir_ls(here::here("data"))
```

```
## /Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/data/p
```

```
dir_ls(here::here("data", "pfiles_sim"))
```

```
## /Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/data/p
## /Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/data/p
## /Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/data/p
## /Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/data/p
## /Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/data/p
## /Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/data/p
## /Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/data/p
```

Limit files

- We really only want the `.CSV`
 - That happens to be the only thing that's in there but that's regularly not the case

```
dir_ls(here::here("data", "pfiles_sim"), glob = "*.csv")
```

[illegible]

Batch load

- Loop through the directories and `import` or `read_csv`

```
files <- dir_ls(  
  here::here("data", "pfiles_sim"),  
  glob = "*.csv"  
)  
batch <- map_dfr(files, read_csv)  
batch
```

```
## # A tibble: 15,945 x 22  
##   Entry   Theta Status Count RawScore      SE Infit Infit_Z Outfit Outfi  
##   <dbl>   <dbl>   <dbl> <dbl>   <dbl>   <dbl> <dbl>   <dbl>   <dbl>   <d  
## 1    123  1.2687     1    36      23 0.3713  0.93   -0.34  0.82   -0  
## 2     88  1.5541     1    36      25 0.3852  0.95   -0.37  0.81   -0  
## 3    105  3.2773     1    36      33 0.6187  0.9    -0.04  1.63    1  
## 4    153  4.4752     1    36      35 1.0234  0.93    0.23  0.35   -0  
## 5    437  2.6655     1    36      31 0.5008  0.92   -0.18  0.88   -0  
## 6    307  5.7137     0    36      36 1.8371  1      0      1      0  
## 7    305  3.7326     1    36      34 0.7408  1.06    0.31  0.86    0  
## 8     42  0.609      1    36      18 0.36    1.55    2.56  1.74    3  
## 9     59 -2.623     1    36       3 1.0344  0.85    0.06  0.17   -0  
## 10   304  5.7137     0    36      36 1.8371  1      0      1      0  
## # ... with 15,935 more rows, and 11 more variables: PointMeasureCorr <dbl>  
## #   ObservMatch <dbl>, ExpectMatch <dbl>, PointMeasureExpected <dbl>, 62RM
```

Problem

- We've lost a lot of info – no way to identify which file is which

Try to fix it!

02:00

Add id

```
batch2 <- map_dfr(files, read_csv, .id = "file")
batch2
```

```
## # A tibble: 15,945 x 23
## # ... with 15,935 more rows, and 23 more variables: file <chr>, Entry <dbl>,
## #   Status <dbl>, Count <dbl>, RawScore <dbl>, SE <dbl>, Infit <dbl>, In
## #   Outfit <dbl>, Outfit_Z <dbl>, Displacement <dbl>, PointMeasureCorr <
## #   ObservMatch <dbl>, ExpectMatch <dbl>, PointMeasureExpected <dbl>, RM
## #   WMLE <dbl>, testeventid <dbl>, ssid <dbl>, asmtprmrydsbltycd <dbl>,
## #   asmtscndrydsbltycd <dbl>
```

Note – the `file` column contains the full path, which is so long it makes no rows print

```
batch2 %>%  
  count(file)
```

```
## # A tibble: 31 x 2  
## # ... with 21 more rows, and 2 more variables: file <chr>, n <int>
```

- Still not terrifically useful. What can we do?

Step 1

- Remove the `here::here` path from string

```
batch2 <- batch2 %>%  
  mutate(  
    file = str_replace_all(  
      file,  
      here::here("data", "pfiles_sim"),  
      ""  
    )  
  )  
  
batch2 %>%  
  count(file)
```

```
## # A tibble: 31 x 2  
##   file                                n  
##   <chr>                             <int>  
## 1 /g11ELApfiles18_sim.csv           453  
## 2 /g11Mathpfiles18_sim.csv          460  
## 3 /g11Rdgpfiles18_sim.csv           453  
## 4 /g11Sciencepfiles18_sim.csv       438  
## 5 /g11Wripfiles18_sim.csv           453  
## 6 /g3ELApfiles18_sim.csv           540
```

Pull out pieces you need

- Regular expressions are most powerful here
 - We haven't talked about them much
- Try RegExplain

Pull grade

- Note – I'm not expecting you to just suddenly be able to do this. This is more for illustration. There's also other ways you could extract the same info

```
batch2 %>%  
  mutate(  
    grade = str_replace_all(  
      file,  
      "/g(\\d?\\d).+",  
      "\\1"  
    )  
  ) %>%  
  select(file, grade)
```

```
## # A tibble: 15,945 x 2  
##   file                                grade  
##   <chr>                             <chr>  
## 1 /g11ELApfiles18_sim.csv 11  
## 2 /g11ELApfiles18_sim.csv 11  
## 3 /g11ELApfiles18_sim.csv 11  
## 4 /g11ELApfiles18_sim.csv 11  
## 5 /g11ELApfiles18_sim.csv 11
```

parse_number

- In this case `parse_number` also works – but note that it would not work to extract the year

```
batch2 %>%  
mutate(grade = parse_number(file)) %>%  
  select(file, grade)
```

```
## # A tibble: 15,945 x 2  
##   file                                grade  
##   <chr>                             <dbl>  
## 1 /g11ELApfiles18_sim.csv           11  
## 2 /g11ELApfiles18_sim.csv           11  
## 3 /g11ELApfiles18_sim.csv           11  
## 4 /g11ELApfiles18_sim.csv           11  
## 5 /g11ELApfiles18_sim.csv           11  
## 6 /g11ELApfiles18_sim.csv           11  
## 7 /g11ELApfiles18_sim.csv           11  
## 8 /g11ELApfiles18_sim.csv           11  
## 9 /g11ELApfiles18_sim.csv           11  
## 10 /g11ELApfiles18_sim.csv          11  
## # ... with 15,935 more rows
```

Extract year

- In this case `parse_number` also works – but note that it would not work to extract the year

```
batch2 %>%  
  mutate(  
    grade = str_replace_all(  
      file, "/g(\\d?\\d).+", "\\1"  
    ),  
    year = str_replace_all(  
      file, ".+files(\\d\\d)_sim.+", "\\1"  
    )  
  ) %>%  
  select(file, grade, year)
```

```
## # A tibble: 15,945 x 3  
##   file                                grade year  
##   <chr>                             <chr> <chr>  
## 1 /g11ELApfiles18_sim.csv 11      18  
## 2 /g11ELApfiles18_sim.csv 11      18  
## 3 /g11ELApfiles18_sim.csv 11      18  
## 4 /g11ELApfiles18_sim.csv 11      18  
## 5 /g11ELApfiles18_sim.csv 11      18
```

Extract Content Area

```
batch2 %>%  
  mutate(grade = str_replace_all(file,  
                                   "/g(\\d?\\d).+",  
                                   "\\1"),  
         year = str_replace_all(file,  
                                 "+files(\\d\\d)_sim.",  
                                 "\\1"),  
         content = str_replace_all(file,  
                                    "/g\\d?\\d(\\d+)pfiles.",  
                                    "\\1")) %>%  
  select(file, grade, year, content)
```

```
## # A tibble: 15,945 x 4  
##   file                                grade year content  
##   <chr>                             <chr> <chr> <chr>  
## 1 /g11ELApfiles18_sim.csv          11    18   ELA  
## 2 /g11ELApfiles18_sim.csv          11    18   ELA  
## 3 /g11ELApfiles18_sim.csv          11    18   ELA  
## 4 /g11ELApfiles18_sim.csv          11    18   ELA  
## 5 /g11ELApfiles18_sim.csv          11    18   ELA  
## 6 /g11ELApfiles18_sim.csv          11    18   ELA  
## 7 /g11ELApfiles18_sim.csv          11    18   ELA  
## 8 /g11ELApfiles18_sim.csv          11    18   ELA  
## 9 /g11ELApfiles18_sim.csv          11    18   ELA
```


Double checks: grade

```
batch2 %>%  
  mutate(grade = str_replace_all(file,  
                                   "/g(\\d?\\d).+",  
                                   "\\1"),  
         year = str_replace_all(file,  
                                   ".+files(\\d\\d)_sim.",  
                                   "\\1"),  
         content = str_replace_all(file,  
                                     "/g\\d?\\d(\\d.+?)pfiles.",  
                                     "\\1")) %>%  
  select(file, grade, year, content) %>%  
  count(grade)
```

```
## # A tibble: 7 x 2  
##   grade      n  
##   <chr> <int>  
## 1 11      2257  
## 2 3       2156  
## 3 4       2341  
## 4 5       2632  
## 5 6       2216  
## 6 7       1962  
## 7 8       2381
```

Double checks: year

```
batch2 %>%
  mutate(grade = str_replace_all(file,
                                   "/g(\\d?\\d).+",
                                   "\\1"),
         year = str_replace_all(file,
                                   ".+files(\\d\\d)_sim.+",
                                   "\\1"),
         content = str_replace_all(file,
                                   "/g\\d?\\d(\\d.+?)pfiles.+",
                                   "\\1")) %>%
  select(file, grade, year, content) %>%
  count(year)
```

```
## # A tibble: 1 x 2
##   year      n
##   <chr> <int>
## 1 18    15945
```

Double checks: content

```
batch2 %>%
  mutate(grade = str_replace_all(file,
                                   "/g(\\d?\\d).+",
                                   "\\1"),
         year = str_replace_all(file,
                                   ".+files(\\d\\d)_sim.+",
                                   "\\1"),
         content = str_replace_all(file,
                                   "/g\\d?\\d(\\d.+?)pfiles.+",
                                   "\\1")) %>%
  select(file, grade, year, content) %>%
  count(content)
```

```
## # A tibble: 5 x 2
##   content      n
##   <chr>    <int>
## 1 ELA      3627
## 2 Math     3629
## 3 Rdg      3627
## 4 Science  1435
## 5 Wri      3627
```

```
d <- batch2 %>%
  mutate(grade = str_replace_all(file, "/g(\\d?\\d).+", "\\1")
         grade = as.integer(grade),
         year = str_replace_all(file, ".+files(\\d\\d)_sim.+",
                                year = as.integer(grade),
                                content = str_replace_all(file, "/g\\d?\\d(.+)pfiles.
  select(-file) %>%
  select(ssid, grade, year, content, testeventid, asmtprmrydsb
         asmtscndrydsbltycd, Entry:WMLE)
```

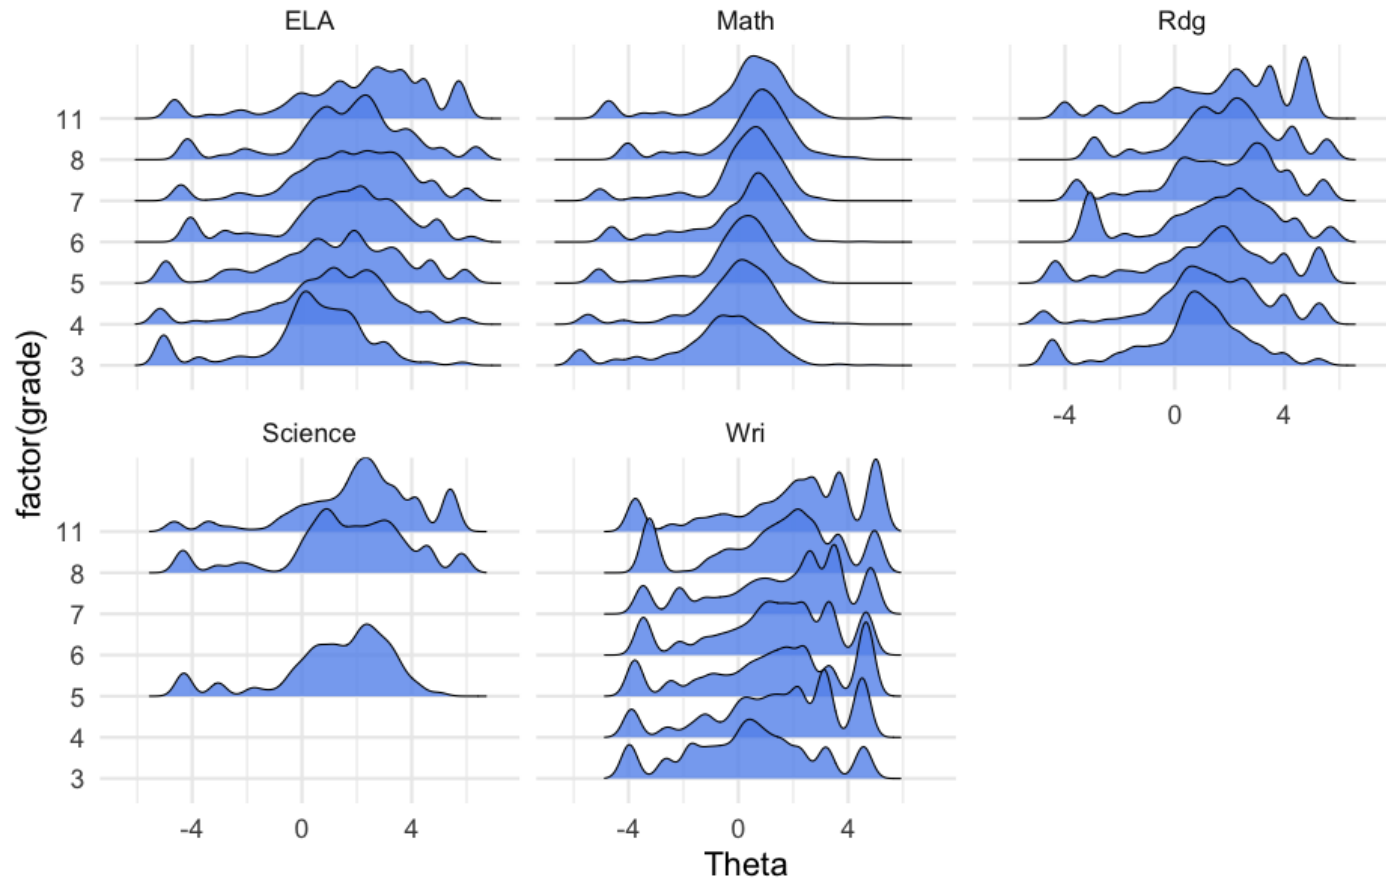
Final product

- In this case, we basically have a tidy data frame already!
- We've reduced our problem from 31 files to a single file

d

```
## # A tibble: 15,945 x 25
##       ssid grade year content testeventid asmtprmrmysdbltycd asmtscndry
##       <dbl> <int> <int> <chr>          <dbl>          <dbl>
##  1  9466908    11    11 ELA             148933            0
##  2  7683685    11    11 ELA             147875           10
##  3  9025693    11    11 ELA             143699           40
##  4 10099824    11    11 ELA             143962           82
##  5 18886078    11    11 ELA             150680           10
##  6 10606750    11    11 ELA             144583           80
##  7 10541306    11    11 ELA             145204           50
##  8  7632967    11    11 ELA             148926           10
##  9  7661118    11    11 ELA             148893           50
## 10 10547177    11    11 ELA             144583           82
## # ... with 15,935 more rows, and 17 more variables: Theta <dbl>, Status <dbl>,
## #   RawScore <dbl>, SE <dbl>, Infit <dbl>, Infit_Z <dbl>, Outfit <dbl>,
## #   Displacement <dbl>, PointMeasureCorr <dbl>, Weight <dbl>, ObservMatch <dbl>,
## #   ExpectMatch <dbl>, PointMeasureExpected <dbl>, RMSR <dbl>, WMLE <dbl>
```

Quick look at distributions



Summary stats

```
d %>%
  group_by(grade, content, asmtprmysbltycd) %>%
  summarize(mean = mean(Theta)) %>%
  pivot_wider(names_from = content,
              values_from = mean)
```

```
## # A tibble: 77 x 7
## # Groups:   grade [7]
##   grade asmtprmysbltycd      ELA      Math      Rdg      Wri
##   <int>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>
## 1      3          0 -0.07361 -1.21055  1.010455  1.612308
## 2      3         10  0.3700416 -0.8182091  0.5184354  0.3206475
## 3      3         20 -0.06335 -1.2514  1.52      -0.5775
## 4      3         40 -1.877683 -3.56365 -1.761667 -0.7514286
## 5      3         50  0.9462857 -0.09186957  0.9791176  1.191481
## 6      3         60  0.840775  1.040375  2.181111  1.067
## 7      3         70 -1.104049 -1.517955 -0.8454839 -1.005625
## 8      3         74  0.996  0.0208375  0.6  1.2925
## 9      3         80 -0.144304 -0.5325596  0.6791667  0.2686301
## 10     3         82  0.3708244 -1.080988  0.5676650  0.3440741
## # ... with 67 more rows
```

Backing up a bit

- What if we wanted only math files?

```
dir_ls(here::here("data", "pfiles_sim"), regexp = "Math")
```

```
## /Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/data/p
## /Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/data/p
## /Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/data/p
## /Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/data/p
## /Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/data/p
## /Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/data/p
## /Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/data/p
```


Only Grade 5

You try

```
g5_paths <- dir_ls(  
  here::here("data", "pfiles_sim"),  
  regexp = "g5"  
)
```

02:00

The rest is the same

```
g5 <- map_dfr(g5_paths, read_csv, .id = "file") %>%
  mutate(
    file = str_replace_all(
      file,
      here::here("data", "pfiles_sim"),
      ""
    )
  )
g5
```

```
## # A tibble: 2,632 x 23
##   file                                Entry  Theta Status Count RawScore      SE Inf
##   <chr>                             <dbl>   <dbl>   <dbl> <dbl>   <dbl>   <dbl> <dbl>
## 1 /g5ELApfiles18_sim.csv           375    3.154     1     36     32 0.551    1.
## 2 /g5ELApfiles18_sim.csv           305    0.3662    1     36     16 0.3894    0.
## 3 /g5ELApfiles18_sim.csv           163   -4.9547   -1     36      0 1.8495    1.
## 4 /g5ELApfiles18_sim.csv           524   -4.9547   -1     36      0 1.8495    1.
## 5 /g5ELApfiles18_sim.csv            81    3.154     1     36     32 0.551    0.
## 6 /g5ELApfiles18_sim.csv           325    1.7156    1     36     25 0.3997    1.
## 7 /g5ELApfiles18_sim.csv           163    1.8786    1     36     26 0.4078    0.
## 8 /g5ELApfiles18_sim.csv           116    5.9323    0     36     36 1.8373    1.
## 9 /g5ELApfiles18_sim.csv           273    1.4052    1     36     23 0.3891    1.
## 10 /g5ELApfiles18_sim.csv          202    1.8786    1     36     26 0.4078    0.
## # ... with 2,622 more rows, and 13 more variables: Outfit_z <dbl>, Displacement_z <dbl>, ...
```

Base equivalents

```
list.files(here::here("data", "pfiles_sim"))
```

```
## [1] "g11ELApfiles18_sim.csv"      "g11Mathpfiles18_sim.csv"    "g11Rdgpfiles18_sim.csv"
## [4] "g11Sciencepfiles18_sim.csv"  "g11Wripfiles18_sim.csv"     "g3ELApfiles18_sim.csv"
## [7] "g3Mathpfiles18_sim.csv"      "g3Rdgpfiles18_sim.csv"     "g3Wripfiles18_sim.csv"
## [10] "g4ELApfiles18_sim.csv"       "g4Mathpfiles18_sim.csv"     "g4Rdgpfiles18_sim.csv"
## [13] "g4Wripfiles18_sim.csv"       "g5ELApfiles18_sim.csv"     "g5Mathpfiles18_sim.csv"
## [16] "g5Rdgpfiles18_sim.csv"       "g5Sciencepfiles18_sim.csv"  "g5Wripfiles18_sim.csv"
## [19] "g6ELApfiles18_sim.csv"       "g6Mathpfiles18_sim.csv"     "g6Rdgpfiles18_sim.csv"
## [22] "g6Wripfiles18_sim.csv"       "g7ELApfiles18_sim.csv"     "g7Mathpfiles18_sim.csv"
## [25] "g7Rdgpfiles18_sim.csv"       "g7Wripfiles18_sim.csv"     "g8ELApfiles18_sim.csv"
## [28] "g8Mathpfiles18_sim.csv"      "g8Rdgpfiles18_sim.csv"     "g8Sciencepfiles18_sim.csv"
## [31] "g8Wripfiles18_sim.csv"
```

Full path

```
list.files(here::here("data", "pfiles_sim"), full.names = TRUE)
```

```
## [1] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/"
## [2] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/"
## [3] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/"
## [4] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/"
## [5] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/"
## [6] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/"
## [7] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/"
## [8] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/"
## [9] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/"
## [10] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/"
## [11] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/"
## [12] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/"
## [13] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/"
## [14] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/"
## [15] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/"
## [16] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/"
## [17] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/"
## [18] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/"
## [19] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/"
## [20] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/"
## [21] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/"
## [22] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/"
## [23] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/"
## [24] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/"
```

Only csvs

```
list.files(here::here("data", "pfiles_sim"),  
          full.names = TRUE,  
          pattern = "*.csv")
```

```
## [1] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/  
## [2] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/  
## [3] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/  
## [4] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/  
## [5] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/  
## [6] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/  
## [7] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/  
## [8] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/  
## [9] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/  
## [10] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/  
## [11] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/  
## [12] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/  
## [13] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/  
## [14] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/  
## [15] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/  
## [16] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/  
## [17] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/  
## [18] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/  
## [19] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/  
## [20] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/  
## [21] "/Users/daniel/Teaching/data_sci_specialization/2020-21/c3-fp-2021/
```

Why not use base?

- We could, but `{fs}` plays a little nicer with `{purrr}`

```
files <- list.files(  
  here::here("data", "pfiles_sim"),  
  pattern = "*.csv"  
)  
batch3 <- map_dfr(files, read_csv, .id = "file")
```

```
## Error: 'g11ELApfiles18_sim.csv' does not exist in current working direct
```

- Need to return full names

```
files
```

```
## [1] "g11ELApfiles18_sim.csv"      "g11Mathpfiles18_sim.csv"    "g11Rdgpfi  
## [4] "g11Sciencepfiles18_sim.csv" "g11Wripfiles18_sim.csv"     "g3ELApfi  
## [7] "g3Mathpfiles18_sim.csv"      "g3Rdgpfiles18_sim.csv"      "g3Wripfi  
## [10] "g4ELApfiles18_sim.csv"       "g4Mathpfiles18_sim.csv"     "g4Rdgpfi  
## [13] "g4Wripfiles18_sim.csv"       "g5ELApfiles18_sim.csv"      "g5Mathpf  
## [16] "g5Rdgpfiles18_sim.csv"       "g5Sciencepfiles18_sim.csv"  "g5Wripfi  
## [19] "g6ELApfiles18_sim.csv"       "g6Mathpfiles18_sim.csv"     "g6Rdgpfi
```

Try again

```
files <- list.files(here::here("data", "pfiles_sim"),
                    pattern = "*.csv",
                    full.names = TRUE)
batch3 <- map_dfr(files, read_csv, .id = "file")
batch3
```

```
## # A tibble: 15,945 x 23
##   file Entry  Theta Status Count RawScore      SE Infit Infit_Z Outfit
##   <chr> <dbl>   <dbl>   <dbl> <dbl>   <dbl>   <dbl> <dbl>   <dbl>   <dbl>
## 1 1      123  1.2687     1     36     23 0.3713  0.93   -0.34  0.82
## 2 1      88  1.5541     1     36     25 0.3852  0.95   -0.37  0.81
## 3 1     105  3.2773     1     36     33 0.6187  0.9    -0.04  1.63
## 4 1     153  4.4752     1     36     35 1.0234  0.93    0.23  0.35
## 5 1     437  2.6655     1     36     31 0.5008  0.92   -0.18  0.88
## 6 1     307  5.7137     0     36     36 1.8371  1       0     1
## 7 1     305  3.7326     1     36     34 0.7408  1.06    0.31  0.86
## 8 1      42  0.609      1     36     18 0.36    1.55    2.56  1.74
## 9 1      59 -2.623      1     36      3 1.0344  0.85    0.06  0.17
## 10 1     304  5.7137     0     36     36 1.8371  1       0     1
## # ... with 15,935 more rows, and 12 more variables: Displacement <dbl>,
## #   PointMeasureCorr <dbl>, Weight <dbl>, ObservMatch <dbl>, ExpectMatch
## #   PointMeasureExpected <dbl>, RMSR <dbl>, WMLE <dbl>, testeventid <dbl>
## #   asmtprmrydsbltycd <dbl>, asmtscndrydsbltycd <dbl>
```

indexes

- The prior example gave us indexes, rather than the file path. Why?

No names

```
names(files)
```

```
## NULL
```

- We **need** the file path! An index isn't nearly as useful.

Base method that works

```
files <- list.files(here::here("data", "pfiles_sim"),
                    pattern = "*.csv",
                    full.names = TRUE)
files <- setNames(files, files)

batch4 <- map_dfr(files, read_csv, .id = "file")
batch4
```

```
## # A tibble: 15,945 x 23
## # ... with 15,935 more rows, and 23 more variables: file <chr>, Entry <dbl>,
## #   Status <dbl>, Count <dbl>, RawScore <dbl>, SE <dbl>, Infit <dbl>, In
## #   Outfit <dbl>, Outfit_Z <dbl>, Displacement <dbl>, PointMeasureCorr <
## #   ObservMatch <dbl>, ExpectMatch <dbl>, PointMeasureExpected <dbl>, RM
## #   WMLE <dbl>, testeventid <dbl>, ssid <dbl>, asmtprmrydsbltycd <dbl>,
## #   asmtscndrydsbltycd <dbl>
```

My recommendation

- If you're working interactively, no reason not to use `{fs}`
- If you are building functions that take generic paths, might be worth considering skipping the dependency

Note

I am **not** saying skip it, but rather that you should **consider** whether it is really needed or not.

List columns

Comparing models

Let's say we wanted to fit/compare a set of models for each content area

1. `lm(Theta ~ asmtprmrydsbltycd)`
2. `lm(Theta ~ asmtprmrydsbltycd +
asmtscndrydsbltycd)`
3. `lm(Theta ~ asmtprmrydsbltycd +
asmtscndrydsbltycd +
asmtprmrydsbltycd:asmtscndrydsbltycd)`

Split the data

The base method we've been using...

```
splt_content <- split(d, d$content)
str(splt_content)
```

```
## List of 5
## $ ELA      : tibble[,25] [3,627 × 25] (S3: tbl_df/tbl/data.frame)
## ..$ ssid      : num [1:3627] 9466908 7683685 9025693 1009982
## ..$ grade     : int [1:3627] 11 11 11 11 11 11 11 11 11 11
## ..$ year      : int [1:3627] 11 11 11 11 11 11 11 11 11 11
## ..$ content   : chr [1:3627] "ELA" "ELA" "ELA" "ELA" ...
## ..$ testeventid : num [1:3627] 148933 147875 143699 143962 150
## ..$ asmtprmrydsblycd : num [1:3627] 0 10 40 82 10 80 50 10 50 82
## ..$ asmtscndrydsblycd : num [1:3627] 0 0 20 0 0 80 0 0 0 0
## ..$ Entry     : num [1:3627] 123 88 105 153 437 307 305 42 5
## ..$ Theta     : num [1:3627] 1.27 1.55 3.28 4.48 2.67 ...
## ..$ Status    : num [1:3627] 1 1 1 1 1 0 1 1 1 0
## ..$ Count     : num [1:3627] 36 36 36 36 36 36 36 36 36 36
## ..$ RawScore  : num [1:3627] 23 25 33 35 31 36 34 18 3 36
## ..$ SE       : num [1:3627] 0.371 0.385 0.619 1.023 0.501
## ..$ Infit     : num [1:3627] 0.93 0.95 0.9 0.93 0.92 1 1.06
## ..$ Infit_Z   : num [1:3627] -0.34 -0.37 -0.04 0.23 -0.18 0
## ..$ Outfit    : num [1:3627] 0.82 0.81 1.63 0.35 0.88 1 0.86
## ..$ Outfit_Z  : num [1:3627] -0.62 -0.56 1.03 -0.16 -0.12 0
```

We could use this method

```
m1 <- map(
  splt_content,
  ~lm(Theta ~ asmtprmrydsbltycd, data = .x)
)
m2 <- map(
  splt_content,
  ~lm(Theta ~ asmtprmrydsbltycd + asmtscndrydsbltycd,
      data = .x)
)
m3 <- map(
  splt_content,
  ~lm(Theta ~ asmtprmrydsbltycd * asmtscndrydsbltycd,
      data = .x)
)
```

- We could then go through and conduct tests to see which model had better fit indices, etc.

Alternative

- Create a data frame with a list column

```
d %>%  
  nest(-content)
```

```
## # A tibble: 5 x 2  
##   content data  
##   <chr>    <list>  
## 1 ELA      <tibble[,24] [3,627 x 24]>  
## 2 Math     <tibble[,24] [3,629 x 24]>  
## 3 Rdg      <tibble[,24] [3,627 x 24]>  
## 4 Science <tibble[,24] [1,435 x 24]>  
## 5 Wri      <tibble[,24] [3,627 x 24]>
```

Add model list column

```
mods <- d %>%
  nest(-content) %>%
  mutate(
    m1 = map(
      data,
      ~lm(Theta ~ asmtprmrydsbltycd,data = .x)
    ),
    m2 = map(
      data,
      ~lm(Theta ~ asmtprmrydsbltycd + asmtscndrydsbltycd,
          data = .x)
    ),
    m3 = map(
      data, ~lm(Theta ~ asmtprmrydsbltycd * asmtscndrydsbltycd
          data = .x)
    )
  )
```


mods

```
## # A tibble: 5 x 5
##   content data                m1      m2      m3
##   <chr>      <list>          <list> <list> <list>
## 1 ELA       <tibble[,24] [3,627 x 24]> <lm>   <lm>   <lm>
## 2 Math      <tibble[,24] [3,629 x 24]> <lm>   <lm>   <lm>
## 3 Rdg       <tibble[,24] [3,627 x 24]> <lm>   <lm>   <lm>
## 4 Science  <tibble[,24] [1,435 x 24]> <lm>   <lm>   <lm>
## 5 Wri       <tibble[,24] [3,627 x 24]> <lm>   <lm>   <lm>
```

Part of the benefit

It's a normal data frame!

```
mods %>%
  pivot_longer(
    m1:m3,
    names_to = "model",
    values_to = "output"
  )
```

```
## # A tibble: 15 x 4
##   content data          model output
##   <chr>   <list>         <chr> <list>
## 1 ELA     <tibble[,24] [3,627 x 24]> m1     <lm>
## 2 ELA     <tibble[,24] [3,627 x 24]> m2     <lm>
## 3 ELA     <tibble[,24] [3,627 x 24]> m3     <lm>
## 4 Math    <tibble[,24] [3,629 x 24]> m1     <lm>
## 5 Math    <tibble[,24] [3,629 x 24]> m2     <lm>
## 6 Math    <tibble[,24] [3,629 x 24]> m3     <lm>
## 7 Rdg     <tibble[,24] [3,627 x 24]> m1     <lm>
## 8 Rdg     <tibble[,24] [3,627 x 24]> m2     <lm>
## 9 Rdg     <tibble[,24] [3,627 x 24]> m3     <lm>
## 10 Science <tibble[,24] [1,435 x 24]> m1     <lm>
## 11 Science <tibble[,24] [1,435 x 24]> m2     <lm>
## 12 Science <tibble[,24] [1,435 x 24]> m3     <lm>
```

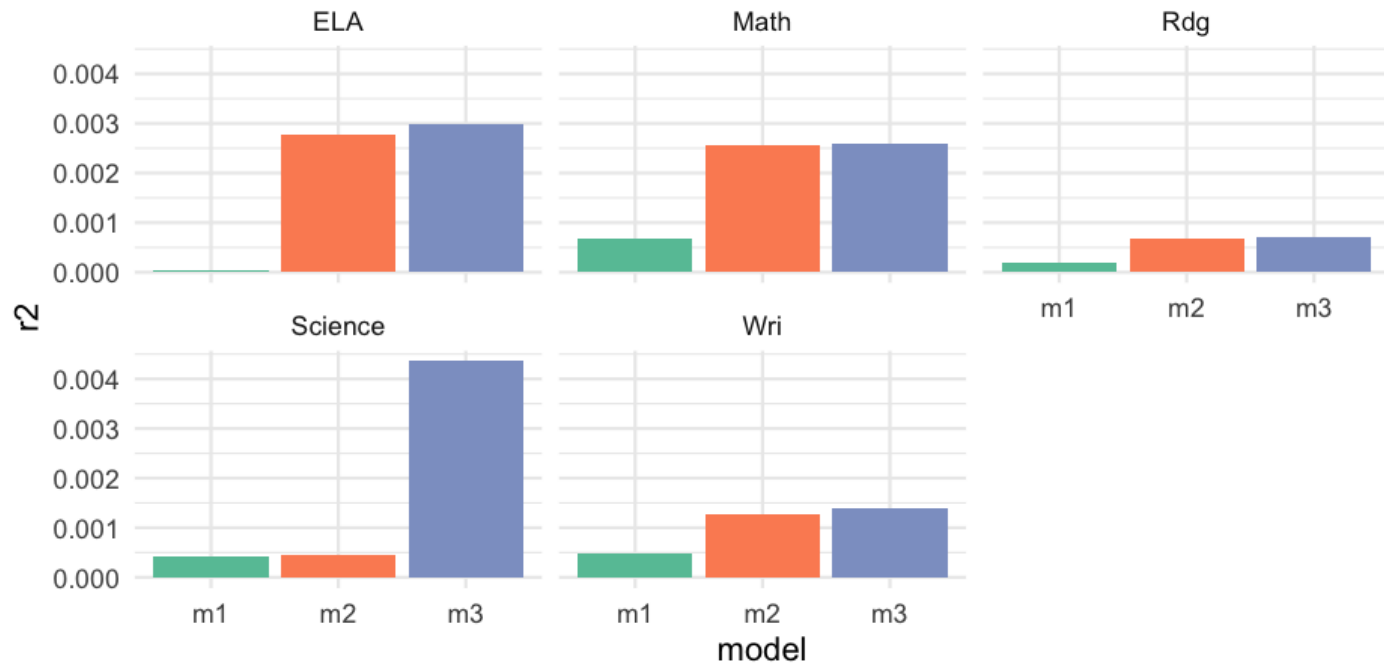
Extract all R^2

```
r2 <- mods %>%
  pivot_longer(
    m1:m3,
    names_to = "model",
    values_to = "output"
  ) %>%
  mutate(r2 = map_dbl(output, ~summary(.x)$r.squared))
r2
```

```
## # A tibble: 15 x 5
##   content data          model output          r2
##   <chr>   <list>         <chr> <list>         <dbl>
## 1 ELA     <tibble[,24] [3,627 x 24]> m1     <lm>      0.00002742625
## 2 ELA     <tibble[,24] [3,627 x 24]> m2     <lm>      0.002784211
## 3 ELA     <tibble[,24] [3,627 x 24]> m3     <lm>      0.002994548
## 4 Math    <tibble[,24] [3,629 x 24]> m1     <lm>      0.0006718361
## 5 Math    <tibble[,24] [3,629 x 24]> m2     <lm>      0.002575408
## 6 Math    <tibble[,24] [3,629 x 24]> m3     <lm>      0.002586228
## 7 Rdg     <tibble[,24] [3,627 x 24]> m1     <lm>      0.0001925962
## 8 Rdg     <tibble[,24] [3,627 x 24]> m2     <lm>      0.0006773540
## 9 Rdg     <tibble[,24] [3,627 x 24]> m3     <lm>      0.0007050212
## 10 Science <tibble[,24] [1,435 x 24]> m1     <lm>      0.0004085780
## 11 Science <tibble[,24] [1,435 x 24]> m2     <lm>      0.0004520424
## 12 Science <tibble[,24] [1,435 x 24]> m3     <lm>      0.004354060
## 13 Wri     <tibble[,24] [3,627 x 24]> m1     <lm>      0.0004902093
```

Plot

```
ggplot(r2, aes(model, r2)) +  
  geom_col(aes(fill = model)) +  
  facet_wrap(~content) +  
  guides(fill = "none") +  
  scale_fill_brewer(palette = "Set2")
```



Summary

- Batch processing is **really** powerful
- Much of the tools we've learned in the past can be applied once we get the data in a more workable format
- List columns are also **really** nice for organization and using our data frame toolkit

