

Agenda

- Review Lab 2
- Introduce list columns
- Contrast:
 - `group_by() %>% nest() %>% mutate() %>% map()` with
 - `nest_by() %>% summarize()`
- In-class midterm (last 20 minutes)

Learning objectives

- Understand list columns and how they relate to `base::split`
- Fluently nest/unnest data frames
- Understand why `tidyr::nest` can be a powerful framework (data frames) and when `tidyr::unnest` can/should be used to move out of nested data frames and into a standard data frame.

Review Lab 2

Setup

Please follow along

First import the data

```
library(tidyverse)
library(fs)
files <- dir_ls(here::here("data", "pfiles_sim"),
               glob = "*.csv")
d <- files %>%
  map_df(read_csv, .id = "file")
```

Parse file data

```
d <- d %>%
  mutate(
    file = str_replace_all(
      file,
      here::here("data", "pfiles_sim"),
      ""
    ),
    grade = str_replace_all(file, "/g(\\d?\\d).+", "\\1"),
    grade = as.integer(grade),
    year = str_replace_all(
      file,
      ".+files(\\d\\d)_sim.+",
      "\\1"
    ),
    year = as.integer(year),
    content = str_replace_all(
      file,
      "/g\\d?\\d(.+)pfiles.+",
      "\\1"
    )
  )
```

Select variables

```
d <- d %>%  
  select(ssid, grade, year, content, testeventid,  
         asmtprmrydsbltycd, asmtscndrydsbltycd, Entry:WMLE)
```

Comparing models

Let's say we wanted to fit/compare a set of models for each content area

1. `lm(Theta ~ asmtprmrydsbltycd)`
2. `lm(Theta ~ asmtprmrydsbltycd +
asmtscndrydsbltycd)`
3. `lm(Theta ~ asmtprmrydsbltycd *
asmtscndrydsbltycd)`

Data pre-processing

- The disability variables are stored as numbers, we need them as factors
- We'll make the names easier in the process

```
d <- d %>%  
  mutate(primary = as.factor(asmtprmrydsbltycd),  
         secondary = as.factor(asmtscndrydsbltycd))
```

If you're interested in what the specific codes refer to, see [here](#).

Split the data

The base method we've been using...

```
splt_content <- split(d, d$content)
str(splt_content)
```

```
## List of 5
## $ ELA      : tibble[,27] [3,627 × 27] (S3: tbl_df/tbl/data.frame)
## ..$ ssid      : num [1:3627] 9466908 7683685 9025693 1009982
## ..$ grade     : int  [1:3627] 11 11 11 11 11 11 11 11 11 11 .
## ..$ year      : int  [1:3627] 18 18 18 18 18 18 18 18 18 18 .
## ..$ content   : chr  [1:3627] "ELA" "ELA" "ELA" "ELA" ...
## ..$ testeventid : num [1:3627] 148933 147875 143699 143962 150
## ..$ asmtprmrydsblycd : num [1:3627] 0 10 40 82 10 80 50 10 50 82 ..
## ..$ asmtscndrydsblycd : num [1:3627] 0 0 20 0 0 80 0 0 0 0 ...
## ..$ Entry     : num [1:3627] 123 88 105 153 437 307 305 42 5
## ..$ Theta     : num [1:3627] 1.27 1.55 3.28 4.48 2.67 ...
## ..$ Status    : num [1:3627] 1 1 1 1 1 0 1 1 1 0 ...
## ..$ Count     : num [1:3627] 36 36 36 36 36 36 36 36 36 36 .
## ..$ RawScore  : num [1:3627] 23 25 33 35 31 36 34 18 3 36 ..
## ..$ SE       : num [1:3627] 0.371 0.385 0.619 1.023 0.501 .
## ..$ Infit     : num [1:3627] 0.93 0.95 0.9 0.93 0.92 1 1.06
## ..$ Infit_Z   : num [1:3627] -0.34 -0.37 -0.04 0.23 -0.18 0
## ..$ Outfit    : num [1:3627] 0.82 0.81 1.63 0.35 0.88 1 0.86
## ..$ Outfit_Z  : num [1:3627] -0.62 -0.56 1.03 -0.16 -0.12 0
## ..$ Displacement : num [1:3627] 0.0018 0.0019 0.0022 0.0023 50.0
```

We could use this method

```
m1 <- map(
  splt_content,
  ~lm(Theta ~ asmtprmrydsbltycd, data = .x)
)

m2 <- map(
  splt_content,
  ~lm(Theta ~ asmtprmrydsbltycd + asmtscndrydsbltycd,
      data = .x)
)

m3 <- map(
  splt_content,
  ~lm(Theta ~ asmtprmrydsbltycd * asmtscndrydsbltycd,
      data = .x)
)
```

- Then conduct tests to see which model fit better, etc.

Alternative

- Create a data frame with a list column

```
by_content <- d %>%  
  group_by(content) %>%  
  nest()  
by_content
```

```
## # A tibble: 5 x 2  
## # Groups:   content [5]  
##   content data  
##   <chr>    <list>  
## 1 ELA      <tibble[,26] [3,627 x 26]>  
## 2 Math     <tibble[,26] [3,629 x 26]>  
## 3 Rdg      <tibble[,26] [3,627 x 26]>  
## 4 Science <tibble[,26] [1,435 x 26]>  
## 5 Wri      <tibble[,26] [3,627 x 26]>
```

What's going on here?

```
str(by_content$data)
```

```
## List of 5
## $ : tibble[,26] [3,627 × 26] (S3: tbl_df/tbl/data.frame)
## ..$ ssid : num [1:3627] 9466908 7683685 9025693 1009982
## ..$ grade : int [1:3627] 11 11 11 11 11 11 11 11 11 11 .
## ..$ year : int [1:3627] 18 18 18 18 18 18 18 18 18 18 .
## ..$ testeventid : num [1:3627] 148933 147875 143699 143962 150
## ..$ asmtprmysblycd : num [1:3627] 0 10 40 82 10 80 50 10 50 82 ..
## ..$ asmtscndrydsblycd : num [1:3627] 0 0 20 0 0 80 0 0 0 0 ...
## ..$ Entry : num [1:3627] 123 88 105 153 437 307 305 42 5
## ..$ Theta : num [1:3627] 1.27 1.55 3.28 4.48 2.67 ...
## ..$ Status : num [1:3627] 1 1 1 1 1 0 1 1 1 0 ...
## ..$ Count : num [1:3627] 36 36 36 36 36 36 36 36 36 36 .
## ..$ RawScore : num [1:3627] 23 25 33 35 31 36 34 18 3 36 ..
## ..$ SE : num [1:3627] 0.371 0.385 0.619 1.023 0.501 .
## ..$ Infit : num [1:3627] 0.93 0.95 0.9 0.93 0.92 1 1.06
## ..$ Infit_Z : num [1:3627] -0.34 -0.37 -0.04 0.23 -0.18 0
## ..$ Outfit : num [1:3627] 0.82 0.81 1.63 0.35 0.88 1 0.86
## ..$ Outfit_Z : num [1:3627] -0.62 -0.56 1.03 -0.16 -0.12 0
## ..$ Displacement : num [1:3627] 0.0018 0.0019 0.0022 0.0023 0.0
## ..$ PointMeasureCorr : num [1:3627] 0.42 0.42 0.3 0.27 0.31 0 0.14
## ..$ Weight : num [1:3627] 1 1 1 1 1 1 1 1 1 1 ...
## ..$ ObservMatch : num [1:3627] 75 80.6 91.7 97.2 86.1 100 94.4
## ..$ ExpectMatch : num [1:3627] 68.3 72 91.7 97.2 86.1 100 94.4
## ..$ PointMeasureExpected: num [1:3627] 0.35 0.33 0.2 0.12 0.25 0 0.17
```

Explore a bit

```
map_dbl(by_content$data, nrow)
```

```
## [1] 3627 3629 3627 1435 3627
```

```
map_dbl(by_content$data, ncol)
```

```
## [1] 26 26 26 26 26
```

```
map_dbl(by_content$data, ~mean(.x$Theta))
```

```
## [1] 1.28001056 -0.06683086 1.37068376 1.57850321 1.26090709
```

It's a data frame!

We can add these summaries if we want

```
by_content %>%  
  mutate(n = map_dbl(data, nrow))
```

```
## # A tibble: 5 x 3  
## # Groups:   content [5]  
##   content data                                n  
##   <chr>    <list>                                <dbl>  
## 1 ELA      <tibble[,26] [3,627 x 26]>    3627  
## 2 Math     <tibble[,26] [3,629 x 26]>    3629  
## 3 Rdg      <tibble[,26] [3,627 x 26]>    3627  
## 4 Science <tibble[,26] [1,435 x 26]>    1435  
## 5 Wri      <tibble[,26] [3,627 x 26]>    3627
```

map_*

- Note on the previous example we used `map_dbl` and we got a vector in return.
- What would happen if we just used `map`?

```
by_content %>%  
  mutate(n = map(data, nrow))
```

```
## # A tibble: 5 x 3  
## # Groups:   content [5]  
##   content data                                n  
##   <chr>    <list>                                <list>  
## 1 ELA     <tibble[,26] [3,627 x 26]> <int [1]>  
## 2 Math    <tibble[,26] [3,629 x 26]> <int [1]>  
## 3 Rdg     <tibble[,26] [3,627 x 26]> <int [1]>  
## 4 Science <tibble[,26] [1,435 x 26]> <int [1]>  
## 5 Wri     <tibble[,26] [3,627 x 26]> <int [1]>
```


Let's fit a model!

```
by_content %>%  
  mutate(m1 = map(data, ~lm(Theta ~ primary, data = .x)))
```

```
## # A tibble: 5 x 3  
## # Groups:   content [5]  
##   content data                                m1  
##   <chr>    <list>                                <list>  
## 1 ELA     <tibble[,26] [3,627 x 26]> <lm>  
## 2 Math    <tibble[,26] [3,629 x 26]> <lm>  
## 3 Rdg     <tibble[,26] [3,627 x 26]> <lm>  
## 4 Science <tibble[,26] [1,435 x 26]> <lm>  
## 5 Wri     <tibble[,26] [3,627 x 26]> <lm>
```

Extract the coefficients

```
by_content %>%  
  mutate(  
    m1 = map(data, ~lm(Theta ~ primary, data = .x)),  
    coefs = map(m1, coef)  
  )
```

```
## # A tibble: 5 x 4  
## # Groups:   content [5]  
##   content data                m1      coefs  
##   <chr>   <list>              <list> <list>  
## 1 ELA    <tibble[,26] [3,627 x 26]> <lm>   <dbl [11]>  
## 2 Math   <tibble[,26] [3,629 x 26]> <lm>   <dbl [12]>  
## 3 Rdg    <tibble[,26] [3,627 x 26]> <lm>   <dbl [11]>  
## 4 Science <tibble[,26] [1,435 x 26]> <lm>   <dbl [12]>  
## 5 Wri    <tibble[,26] [3,627 x 26]> <lm>   <dbl [12]>
```

Challenge

- Continue with the above, but output a data frame with three columns: **content**, **intercept**, and **TBI** (which is code 74).
- In other words, output the mean score for students who were coded as not having a disability (code 0), along with students coded as having TBI.

04:00

```
by_content %>%
  mutate(
    m1 = map(data, ~lm(Theta ~ primary, data = .x)),
    coefs = map(m1, coef),
    no_disab = map_dbl(coefs, 1),
    tbi = no_disab + map_dbl(coefs, "primary74")
  ) %>%
  select(content, no_disab, tbi)
```

```
## # A tibble: 5 x 3
## # Groups:   content [5]
##   content    no_disab      tbi
##   <chr>      <dbl>    <dbl>
## 1 ELA        0.9322336  0.1674462
## 2 Math      -0.1587907  0.1910821
## 3 Rdg        1.363101   1.629048
## 4 Science    1.491319   2.790971
## 5 Wri        1.571441   1.167429
```

Note – I wouldn't have necessarily expected you to add `no_disab` to the TBI coefficient.

Compare models

- Back to our original task – fit all three models

You try first

1. `lm(Theta ~ primary)`
2. `lm(Theta ~ primary + secondary)`
3. `lm(Theta ~ primary + secondary +
primary:secondary)`

04:00

Model fits

```
mods <- by_content %>%
  mutate(
    m1 = map(data, ~lm(Theta ~ primary, data = .x)),
    m2 = map(data, ~lm(Theta ~ primary + secondary, data = .x)),
    m3 = map(data, ~lm(Theta ~ primary * secondary, data = .x))
  )
mods
```

```
## # A tibble: 5 x 5
## # Groups:   content [5]
##   content data                m1      m2      m3
##   <chr>   <list>              <list> <list> <list>
## 1 ELA     <tibble[,26] [3,627 x 26]> <lm>   <lm>   <lm>
## 2 Math    <tibble[,26] [3,629 x 26]> <lm>   <lm>   <lm>
## 3 Rdg     <tibble[,26] [3,627 x 26]> <lm>   <lm>   <lm>
## 4 Science <tibble[,26] [1,435 x 26]> <lm>   <lm>   <lm>
## 5 Wri     <tibble[,26] [3,627 x 26]> <lm>   <lm>   <lm>
```

Brief foray into parallel iterations

The `stats::anova` function can compare the fit of two models

Pop Quiz

How would we extract just ELA model 1 and 2?

```
mods$m1[[1]]
```

```
##  
## Call:  
## lm(formula = Theta ~ primary, data = data, weights = wts)  
##  
## Coefficients:  
## (Intercept)      primary10  
##      0.93223      0.38570
```

```
mods$m2[[1]]
```

```
##  
## Call:  
## lm(formula = Theta ~ primary + second, data = data, weights = wts)  
##  
## Coefficients:  
## (Intercept)      primary10      primary20      primary40      primary50      primary20  
##      -0.03168      1.00143      0.42185      0.42185      0.28070
```

Which fits better?

```
compare <- anova(mods$m1[[1]], mods$m2[[1]])  
compare
```

```
## Analysis of Variance Table  
##  
## Model 1: Theta ~ primary  
## Model 2: Theta ~ primary + secondary  
##   Res.Df    RSS Df Sum of Sq      F    Pr(>F)  
## 1     3616 20905  
## 2     3605 20100 11     804.26 13.113 < 2.2e-16 ***  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```


map2

- Works the same as `map` but iterates over two vectors concurrently
- Let's compare model 1 and 2

```
mods %>%  
  mutate(comp12 = map2(m1, m2, anova))
```

```
## # A tibble: 5 x 6  
## # Groups:   content [5]  
##   content data          m1      m2      m3      comp12  
##   <chr>   <list>         <list> <list> <list> <list>  
## 1 ELA    <tibble[,26] [3,627 x 26]> <lm>   <lm>   <lm>   <anova[,6] [2  
## 2 Math   <tibble[,26] [3,629 x 26]> <lm>   <lm>   <lm>   <anova[,6] [2  
## 3 Rdg    <tibble[,26] [3,627 x 26]> <lm>   <lm>   <lm>   <anova[,6] [2  
## 4 Science <tibble[,26] [1,435 x 26]> <lm>   <lm>   <lm>   <anova[,6] [2  
## 5 Wri    <tibble[,26] [3,627 x 26]> <lm>   <lm>   <lm>   <anova[,6] [2
```

Perhaps not terrifically helpful

Back to our **anova** object

- Can we pull out useful things?

```
str(compare)
```

```
## Classes 'anova' and 'data.frame':    2 obs. of  6 variables:
##  $ Res.Df    : num  3616 3605
##  $ RSS       : num  20905 20100
##  $ Df        : num  NA 11
##  $ Sum of Sq: num  NA 804
##  $ F         : num  NA 13.1
##  $ Pr(>F)    : num  NA 7.66e-25
##  - attr(*, "heading")= chr [1:2] "Analysis of Variance Table\n" "Model 1"
```

Try pulling out the *p* value

Extract *p* value

- *Note – I'd recommend looking at more than just a p-value, but I do think this is useful for a quick glance*

```
compare$`Pr(>F)`
```

```
## [1] NA 7.663566e-25
```

```
compare[["Pr(>F)"]]
```

```
## [1] NA 7.663566e-25
```

```
compare$`Pr(>F)`[2]
```

```
## [1] 7.663566e-25
```

```
compare[["Pr(>F)"]][2]
```

```
## [1] 7.663566e-25
```

All p-values

Note – this is probably the most compact syntax, but that doesn't mean it's the most clear

```
mods %>%  
  mutate(comp12 = map2(m1, m2, anova),  
         p12 = map_dbl(comp12, list("Pr(>F)", 2)))
```

```
## # A tibble: 5 x 7  
## # Groups:   content [5]  
##   content data          m1      m2      m3      comp12  
##   <chr>   <list>         <list> <list> <list> <list>  
## 1 ELA    <tibble[,26] [3,627 x 26]> <lm>   <lm>   <lm>   <anova[,6] [2  
## 2 Math   <tibble[,26] [3,629 x 26]> <lm>   <lm>   <lm>   <anova[,6] [2  
## 3 Rdg    <tibble[,26] [3,627 x 26]> <lm>   <lm>   <lm>   <anova[,6] [2  
## 4 Science <tibble[,26] [1,435 x 26]> <lm>   <lm>   <lm>   <anova[,6] [2  
## 5 Wri    <tibble[,26] [3,627 x 26]> <lm>   <lm>   <lm>   <anova[,6] [2
```

Slight alternative

- Write a function that pulls the p-value from model comparison objects

```
extract_p <- function(anova_ob) {  
  anova_ob[["Pr(>F)"]][2]  
}
```

- Loop this function through the anova objects

```

mods %>%
  mutate(comp12 = map2(m1, m2, anova),
         p12 = map_dbl(comp12, extract_p))

```

```

## # A tibble: 5 x 7
## # Groups:   content [5]
##   content data          m1      m2      m3      comp12
##   <chr>   <list>      <list> <list> <list> <list>
## 1 ELA     <tibble[,26] [3,627 x 26]> <lm>   <lm>   <lm>   <anova[,6] [2
## 2 Math    <tibble[,26] [3,629 x 26]> <lm>   <lm>   <lm>   <anova[,6] [2
## 3 Rdg     <tibble[,26] [3,627 x 26]> <lm>   <lm>   <lm>   <anova[,6] [2
## 4 Science <tibble[,26] [1,435 x 26]> <lm>   <lm>   <lm>   <anova[,6] [2
## 5 Wri     <tibble[,26] [3,627 x 26]> <lm>   <lm>   <lm>   <anova[,6] [2

```

An alternative

Conducting operations by row

Operations by row

The `dplyr::rowwise()` function fundamentally changes the way a `tibble()` behaves

```
df <- tibble(name = c("Me", "You"), x = 1:2, y = 3:4, z = 5:6)
```

```
df %>%  
  mutate(m = mean(c(x, y, z)))
```

```
## # A tibble: 2 x 5  
##   name      x      y      z      m  
##   <chr> <int> <int> <int> <dbl>  
## 1 Me         1         3         5     3.5  
## 2 You         2         4         6     3.5
```

```
df %>%  
  rowwise() %>%  
  mutate(m = mean(c(x, y, z)))
```

```
## # A tibble: 2 x 5  
## # Rowwise:  
##   name      x      y      z      m  
##   <chr> <int> <int> <int> <dbl>  
## 1 Me         1         3         5     3  
## 2 You         2         4         6     4
```


Add a group & summarize

```
df %>%  
  rowwise(name) %>%  
  summarize(m = mean(c(x, y, z)))
```

```
## # A tibble: 2 x 2  
## # Groups:   name [2]  
##   name      m  
##   <chr> <dbl>  
## 1 Me      3  
## 2 You      4
```

List columns

If you apply rowwise operation with a list column, you don't have to loop

```
df <- tibble(x = list(1, 2:3, 4:6))
```

```
df %>%  
  mutate(  
    l = map_int(x, length)  
  )
```

```
## # A tibble: 3 x 2  
##   x          1  
##   <list>    <int>  
## 1 <dbl [1]>    1  
## 2 <int [2]>    2  
## 3 <int [3]>    3
```

```
df %>%  
  rowwise() %>%  
  mutate(l = length(x))
```

```
## # A tibble: 3 x 2  
## # Rowwise:  
##   x          1  
##   <list>    <int>  
## 1 <dbl [1]>    1  
## 2 <int [2]>    2  
## 3 <int [3]>    3
```

Creating list columns

You can use the `dplyr::nest_by()` function to create a list column for each group, *and* convert it to a rowwise data frame.

```
d %>%  
  nest_by(content)
```

```
## # A tibble: 5 x 2  
## # Rowwise:   content  
##   content      data  
##   <chr>    <list<tibble[,26]>>  
## 1 ELA      [3,627 x 26]  
## 2 Math     [3,629 x 26]  
## 3 Rdg      [3,627 x 26]  
## 4 Science  [1,435 x 26]  
## 5 Wri      [3,627 x 26]
```

Challenge

Given what we just learned, can you fit a model of the form **Theta ~ primary** to each content area (i.e., *not* using **{purrr}**)?

Wrap it in **list()** (should suggest this in the error reporting if you don't)

```
d %>%
  nest_by(content) %>%
  mutate(m1 = list(lm(Theta ~ primary, data = data)))
```

```
## # A tibble: 5 x 3
## # Rowwise:   content
##   content                data m1
##   <chr>      <list<tibble[,26]>> <list>
## 1 ELA          [3,627 x 26] <lm>
## 2 Math          [3,629 x 26] <lm>
## 3 Rdg           [3,627 x 26] <lm>
## 4 Science       [1,435 x 26] <lm>
## 5 Wri           [3,627 x 26] <lm>
```

02:00

Challenge 2

Can you extend it further and extract the coefficients with **coef**? What about creating a new column that has the intercept values?

```
d %>%  
  nest_by(content) %>%  
  mutate(m1 = list(lm(Theta ~ primary, data = data)),  
         coefs = list(coef(m1)))
```

```
## # A tibble: 5 x 4  
## # Rowwise:   content  
##   content      data m1      coefs  
##   <chr>    <list<tibble[,26]>> <list> <list>  
## 1 ELA      [3,627 x 26] <lm>    <dbl [11]>  
## 2 Math     [3,629 x 26] <lm>    <dbl [12]>  
## 3 Rdg      [3,627 x 26] <lm>    <dbl [11]>  
## 4 Science [1,435 x 26] <lm>    <dbl [12]>  
## 5 Wri      [3,627 x 26] <lm>    <dbl [12]>
```

02:00

Return atomic vectors

```
d %>%  
  nest_by(content) %>%  
  mutate(m1 = list(lm(Theta ~ primary, data = data)),  
         intercept = coef(m1)[1])
```

```
## # A tibble: 5 x 4  
## # Rowwise:  content  
##   content          data m1      intercept  
##   <chr>    <list<tibble[,26]>> <list>      <dbl>  
## 1 ELA      [3,627 x 26] <lm>      0.9322336  
## 2 Math      [3,629 x 26] <lm>     -0.1587907  
## 3 Rdg       [3,627 x 26] <lm>      1.363101  
## 4 Science   [1,435 x 26] <lm>      1.491319  
## 5 Wri       [3,627 x 26] <lm>      1.571441
```

Fit all models

The below gets us the same results we got before

```
mods2 <- d %>%
  nest_by(content) %>%
  mutate(
    m1 = list(lm(Theta ~ primary, data = data)),
    m2 = list(lm(Theta ~ primary + secondary, data = data)),
    m3 = list(lm(Theta ~ primary * secondary, data = data))
  )
mods2
```

```
## # A tibble: 5 x 5
## # Rowwise:  content
##   content      data m1      m2      m3
##   <chr>    <list<tibble[,26]>> <list> <list> <list>
## 1 ELA      [3,627 x 26] <lm>   <lm>   <lm>
## 2 Math      [3,629 x 26] <lm>   <lm>   <lm>
## 3 Rdg       [3,627 x 26] <lm>   <lm>   <lm>
## 4 Science  [1,435 x 26] <lm>   <lm>   <lm>
## 5 Wri       [3,627 x 26] <lm>   <lm>   <lm>
```

Look at all R^2

It's a normal data frame!

```
mods %>%
  pivot_longer(
    m1:m3,
    names_to = "model",
    values_to = "output"
  )
```

```
## # A tibble: 15 x 4
## # Groups:   content [5]
##   content data                                model output
##   <chr>   <list>                                <chr> <list>
## 1 ELA    <tibble[,26] [3,627 x 26]> m1    <lm>
## 2 ELA    <tibble[,26] [3,627 x 26]> m2    <lm>
## 3 ELA    <tibble[,26] [3,627 x 26]> m3    <lm>
## 4 Math   <tibble[,26] [3,629 x 26]> m1    <lm>
## 5 Math   <tibble[,26] [3,629 x 26]> m2    <lm>
## 6 Math   <tibble[,26] [3,629 x 26]> m3    <lm>
## 7 Rdg    <tibble[,26] [3,627 x 26]> m1    <lm>
## 8 Rdg    <tibble[,26] [3,627 x 26]> m2    <lm>
## 9 Rdg    <tibble[,26] [3,627 x 26]> m3    <lm>
## 10 Science <tibble[,26] [1,435 x 26]> m1    <lm>
## 11 Science <tibble[,26] [1,435 x 26]> m2    <lm>
```


Extract all R^2

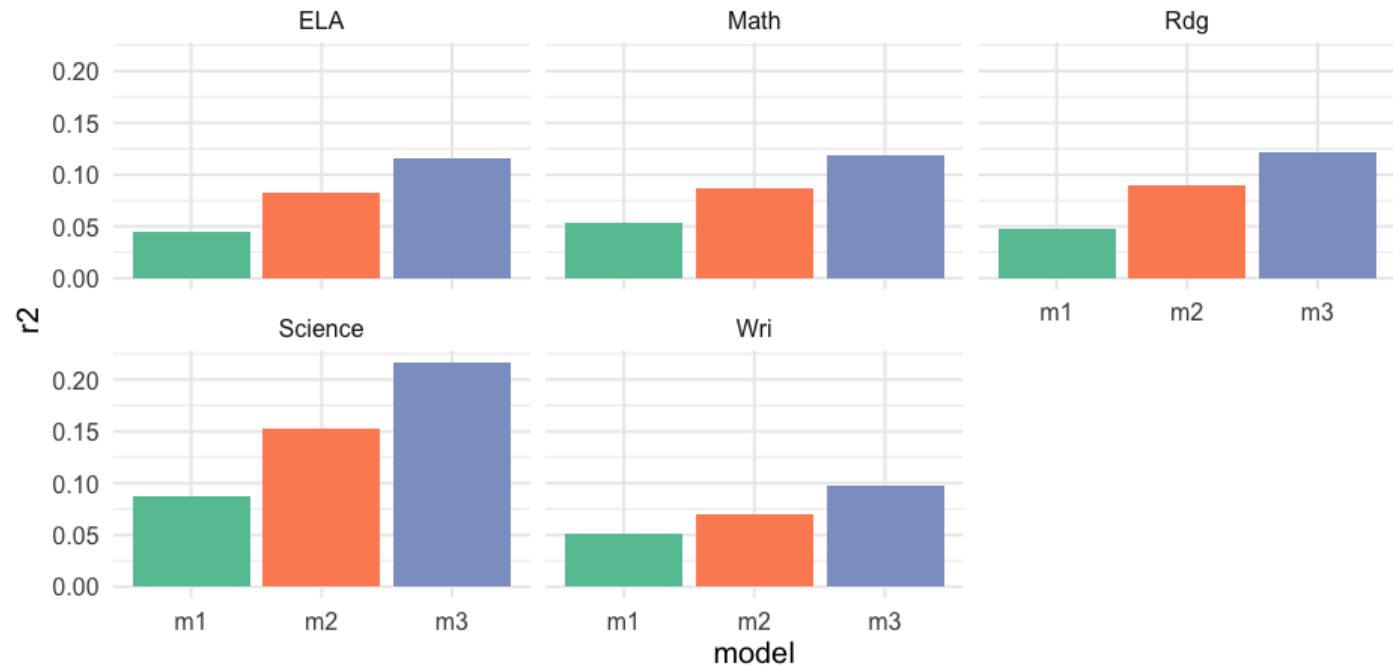
Note – might want to write a function here again

```
r2 <- mods %>%
  pivot_longer(
    m1:m3,
    names_to = "model",
    values_to = "output"
  ) %>%
  mutate(r2 = map_dbl(output, ~summary(.x)$r.squared))
r2
```

```
## # A tibble: 15 x 5
## # Groups:   content [5]
##   content data                model output          r2
##   <chr>   <list>              <chr> <list>        <dbl>
## 1 ELA     <tibble[,26] [3,627 x 26]> m1     <lm>      0.04517421
## 2 ELA     <tibble[,26] [3,627 x 26]> m2     <lm>      0.08190917
## 3 ELA     <tibble[,26] [3,627 x 26]> m3     <lm>      0.1161187
## 4 Math    <tibble[,26] [3,629 x 26]> m1     <lm>      0.05326550
## 5 Math    <tibble[,26] [3,629 x 26]> m2     <lm>      0.08675264
## 6 Math    <tibble[,26] [3,629 x 26]> m3     <lm>      0.1185931
## 7 Rdg     <tibble[,26] [3,627 x 26]> m1     <lm>      0.04805713
## 8 Rdg     <tibble[,26] [3,627 x 26]> m2     <lm>      0.08926212
## 9 Rdg     <tibble[,26] [3,627 x 26]> m3     <lm>      0.1217497
```

Plot

```
ggplot(r2, aes(model, r2)) +  
  geom_col(aes(fill = model)) +  
  facet_wrap(~content) +  
  guides(fill = "none") +  
  scale_fill_brewer(palette = "Set2")
```



Unnesting

- Sometimes you just want to `unnest`
- Imagine we want to plot the coefficients by model... how?
- `broom::tidy()` => `tidyr::unnest()`

Tidy

```
mods %>%
  pivot_longer(
    m1:m3,
    names_to = "model",
    values_to = "output"
  ) %>%
  mutate(tidied = map(output, broom::tidy))
```

```
## # A tibble: 15 x 5
## # Groups:   content [5]
##   content data                                model output tidied
##   <chr>   <list>                                <chr> <list> <list>
## 1 ELA     <tibble[,26] [3,627 x 26]> m1     <lm>   <tibble[,5] [11 x 5]>
## 2 ELA     <tibble[,26] [3,627 x 26]> m2     <lm>   <tibble[,5] [22 x 5]>
## 3 ELA     <tibble[,26] [3,627 x 26]> m3     <lm>   <tibble[,5] [132 x 5]>
## 4 Math    <tibble[,26] [3,629 x 26]> m1     <lm>   <tibble[,5] [12 x 5]>
## 5 Math    <tibble[,26] [3,629 x 26]> m2     <lm>   <tibble[,5] [23 x 5]>
## 6 Math    <tibble[,26] [3,629 x 26]> m3     <lm>   <tibble[,5] [144 x 5]>
## 7 Rdg     <tibble[,26] [3,627 x 26]> m1     <lm>   <tibble[,5] [11 x 5]>
## 8 Rdg     <tibble[,26] [3,627 x 26]> m2     <lm>   <tibble[,5] [22 x 5]>
## 9 Rdg     <tibble[,26] [3,627 x 26]> m3     <lm>   <tibble[,5] [132 x 5]>
## 10 Science <tibble[,26] [1,435 x 26]> m1     <lm>   <tibble[,5] [12 x 5]>
## 11 Science <tibble[,26] [1,435 x 26]> m2     <lm>   <tibble[,5] [22 x 5]>
## 12 Science <tibble[,26] [1,435 x 26]> m3     <lm>   <tibble[,5] [132 x 5]>
## 13 Wri     <tibble[,26] [3,627 x 26]> m1     <lm>   <tibble[,5] [12 x 5]>
```

Equivalently

```
mods %>%
  pivot_longer(
    m1:m3,
    names_to = "model",
    values_to = "output"
  ) %>%
  rowwise() %>%
  mutate(tidied = list(broom::tidy(output)))
```

```
## # A tibble: 15 x 5
## # Rowwise:  content
##   content data          model output tidied
##   <chr>   <list>         <chr> <list> <list>
## 1 ELA    <tibble[,26] [3,627 x 26]> m1    <lm>   <tibble[,5] [11 x 5]>
## 2 ELA    <tibble[,26] [3,627 x 26]> m2    <lm>   <tibble[,5] [22 x 5]>
## 3 ELA    <tibble[,26] [3,627 x 26]> m3    <lm>   <tibble[,5] [132 x 5]>
## 4 Math   <tibble[,26] [3,629 x 26]> m1    <lm>   <tibble[,5] [12 x 5]>
## 5 Math   <tibble[,26] [3,629 x 26]> m2    <lm>   <tibble[,5] [23 x 5]>
## 6 Math   <tibble[,26] [3,629 x 26]> m3    <lm>   <tibble[,5] [144 x 5]>
## 7 Rdg    <tibble[,26] [3,627 x 26]> m1    <lm>   <tibble[,5] [11 x 5]>
## 8 Rdg    <tibble[,26] [3,627 x 26]> m2    <lm>   <tibble[,5] [22 x 5]>
## 9 Rdg    <tibble[,26] [3,627 x 26]> m3    <lm>   <tibble[,5] [132 x 5]>
## 10 Science <tibble[,26] [1,435 x 26]> m1    <lm>   <tibble[,5] [12 x 5]>
## 11 Science <tibble[,26] [1,435 x 26]> m2    <lm>   <tibble[,5] [22 x 5]>
## 12 Science <tibble[,26] [1,435 x 26]> m3    <lm>   <tibble[,5] [132 x 5]
```

Select and unnest

```
tidied <- mods %>%  
  gather(model, output, m1:m3) %>%  
  mutate(tidied = map(output, broom::tidy)) %>%  
  select(content, model, tidied) %>%  
  unnest(tidied)  
tidied
```


```
## # A tibble: 841 x 7  
## # Groups:   content [5]  
##   content model term          estimate std.error    statistic    p.val  
##   <chr>    <chr> <chr>          <dbl>      <dbl>      <dbl>      <dbl>  
## 1 ELA      m1      (Intercept)  0.9322336  0.2150561  4.334839   1.498396e-05  
## 2 ELA      m1      primary10    0.3856986  0.2242965  1.719593   8.559207e-02  
## 3 ELA      m1      primary20   -0.03167527 0.7266436 -0.04359120 9.652327e-02  
## 4 ELA      m1      primary40   -1.844343    0.5559031 -3.317741   9.164595e-03  
## 5 ELA      m1      primary50    1.173722    0.2890447  4.060694   4.996391e-05  
## 6 ELA      m1      primary60    0.7762539   0.3866313  2.007737   4.474555e-02  
## 7 ELA      m1      primary70   -1.830257    0.3086128 -5.930595   3.301860e-06  
## 8 ELA      m1      primary74   -0.7647874   0.5182670 -1.475663   1.401215e-01  
## 9 ELA      m1      primary80    0.4676481   0.2428640  1.925556   5.423822e-02  
## 10 ELA     m1      primary82    0.3382547   0.2267600  1.491686   1.358687e-01  
## # ... with 831 more rows
```

Plot

Lets look how the primary coefficients change

```
to_plot <- names(coef(mods$m1[[1]]))

tidied %>%
  filter(term %in% to_plot) %>%
  ggplot(aes(estimate, term, color = model)) +
  geom_point() +
  scale_color_brewer(palette = "Set2") +
  facet_wrap(~content)
```



Last bit

- We've kind of been running the wrong models this whole time
- We forgot about grade!
- No problem, just change the grouping factor

By grade

```
by_grade_content <- d %>%  
  group_by(content, grade) %>%  
  nest()  
by_grade_content
```

```
## # A tibble: 31 x 3  
## # Groups:   grade, content [31]  
##   grade content data  
##   <int> <chr>   <list>  
## 1     11 ELA     <tibble[,25] [453 x 25]>  
## 2     11 Math    <tibble[,25] [460 x 25]>  
## 3     11 Rdg     <tibble[,25] [453 x 25]>  
## 4     11 Science <tibble[,25] [438 x 25]>  
## 5     11 Wri     <tibble[,25] [453 x 25]>  
## 6      3 ELA     <tibble[,25] [540 x 25]>  
## 7      3 Math    <tibble[,25] [536 x 25]>  
## 8      3 Rdg     <tibble[,25] [540 x 25]>  
## 9      3 Wri     <tibble[,25] [540 x 25]>  
## 10     4 ELA     <tibble[,25] [585 x 25]>  
## # ... with 21 more rows
```

Fit models

```
mods_grade <- by_grade_content %>%
  mutate(
    m1 = map(data, ~lm(Theta ~ primary, data = .x)),
    m2 = map(data, ~lm(Theta ~ primary + secondary,
                      data = .x)),
    m3 = map(data, ~lm(Theta ~ primary * secondary,
                      data = .x))
  )
mods_grade
```

```
## # A tibble: 31 x 6
## # Groups:   grade, content [31]
##   grade content data          m1      m2      m3
##   <int> <chr>   <list>    <list> <list> <list>
## 1     11 ELA    <tibble[,25] [453 x 25]> <lm>   <lm>   <lm>
## 2     11 Math   <tibble[,25] [460 x 25]> <lm>   <lm>   <lm>
## 3     11 Rdg    <tibble[,25] [453 x 25]> <lm>   <lm>   <lm>
## 4     11 Science <tibble[,25] [438 x 25]> <lm>   <lm>   <lm>
## 5     11 Wri    <tibble[,25] [453 x 25]> <lm>   <lm>   <lm>
## 6      3 ELA    <tibble[,25] [540 x 25]> <lm>   <lm>   <lm>
## 7      3 Math   <tibble[,25] [536 x 25]> <lm>   <lm>   <lm>
## 8      3 Rdg    <tibble[,25] [540 x 25]> <lm>   <lm>   <lm>
## 9      3 Wri    <tibble[,25] [540 x 25]> <lm>   <lm>   <lm>
## 10     4 ELA    <tibble[,25] [585 x 25]> <lm>   <lm>   <lm>
## # ... with 21 more rows
```

Look at R^2

```
mods_grade %>%
  pivot_longer(
    m1:m3,
    names_to = "model",
    values_to = "output"
  ) %>%
  mutate(r2 = map_dbl(output, ~summary(.x)$r.squared))
```

```
## # A tibble: 93 x 6
## # Groups:   grade, content [31]
##   grade content data          model output      r2
##   <int> <chr>   <list>         <chr> <list>    <dbl>
## 1     11 ELA    <tibble[,25] [453 x 25]> m1    <lm>     0.03353818
## 2     11 ELA    <tibble[,25] [453 x 25]> m2    <lm>     0.1084394
## 3     11 ELA    <tibble[,25] [453 x 25]> m3    <lm>     0.1536891
## 4     11 Math   <tibble[,25] [460 x 25]> m1    <lm>     0.1886003
## 5     11 Math   <tibble[,25] [460 x 25]> m2    <lm>     0.3161226
## 6     11 Math   <tibble[,25] [460 x 25]> m3    <lm>     0.4046634
## 7     11 Rdg    <tibble[,25] [453 x 25]> m1    <lm>     0.02066316
## 8     11 Rdg    <tibble[,25] [453 x 25]> m2    <lm>     0.1820512
## 9     11 Rdg    <tibble[,25] [453 x 25]> m3    <lm>     0.2337721
## 10    11 Science <tibble[,25] [438 x 25]> m1    <lm>     0.1259080
## # ... with 83 more rows
```

Plot

```
mods_grade %>%  
  pivot_longer(  
    m1:m3,  
    names_to = "model",  
    values_to = "output"  
  ) %>%  
  mutate(r2 = map_dbl(output, ~summary(.x)$r.squared)) %>%  
  ggplot(aes(model, r2)) +  
  geom_col(aes(fill = model)) +  
  facet_grid(grade ~ content) +  
  guides(fill = "none") +  
  scale_fill_brewer(palette = "Set2")
```


Summary

- List columns are really powerful and really flexible
- Also help you stay organized
- You can approach the problem either with `{purrr}` or `dplyr::rowwise()`.
 - **Important:** If you use `rowwise()`, remember to `ungroup()` when you want it to go back to being a normal data frame
 - I'm asking you to learn both – the row-wise approach might be a bit easier but is a little less general (only works with data frames)

Questions?

In-class

Midterm

Next time: Parallel iterations