

Looping Variants

Daniel Anderson

Week 5, Class 2

Agenda

- `walk()` and friends
- `modify()`
- `safely()`
- `reduce()`

Learning Objectives

- Know when to apply `walk` instead of `map`, and why it may be useful
- Understand the parallels and differences between `map` and `modify`
- Diagnose errors with `safely` and understand other situations where it may be helpful
- Collapsing/reducing lists with `purrr::reduce()` or `base::Reduce()`

Setup

Let's go back to our plotting example from last class.

First we'll load our libraries

```
library(tidyverse)  
library(fivethirtyeight)  
library(glue)  
library(english)
```

Prep the data

```
pulitzer <- pulitzer %>%  
  select(newspaper, starts_with("num")) %>%  
  pivot_longer(  
    -newspaper,  
    names_to = "year_range",  
    values_to = "n",  
    names_prefix = "num_finals"  
  ) %>%  
  mutate(year_range = str_replace_all(year_range, "_", "-")) %>%  
  filter(year_range != "1990-2014") %>%  
  group_by(newspaper) %>%  
  mutate(  
    tot = sum(n),  
    label = glue(  
      "{str_to_title(as.english(tot))} Total Pulitzer Awards"  
    )  
  )  
)
```

Produce plots

```
final_plots <- pulitzer %>%
  group_by(newspaper, label) %>%
  nest() %>%
  mutate(plots = pmap(list(newspaper, label, data), ~{
    ggplot(..3, aes(n, year_range)) +
      geom_col(aes(fill = n)) +
      scale_fill_distiller(type = "seq",
                           limits = c(0, max(pulitzer$n)),
                           palette = "BuPu",
                           direction = 1) +
      scale_x_continuous(limits = c(0, max(pulitzer$n)),
                         expand = c(0, 0)) +
      guides(fill = "none") +
      labs(title = glue("Pulitzer Prize winners: {..1}"),
           x = "Total number of winners",
           y = "",
           caption = ..2)
  })
)
```

Saving

- We saw last time that we could use `nest_by()`
 - Required a bit of awkwardness with adding the paths to the data frame
 - Instead, we'll do it again but with the `walk()` family

Why `walk()` for saving instead of `map()`?

Walk is an alternative to map that you use when you want to call a function for its side effects, rather than for its return value. You typically do this because you want to render output to the screen or save files to disk – the important thing is the action, not the return value.

More practical

If you use `walk()`, nothing will get printed to the screen.
This is particularly helpful for RMarkdown files.

Example

Please do the following

- Create a new RMarkdown document
- Paste the code you have for creating the plots in a code chunk there (along with the library loading, data cleaning, etc.)

02:00

Create a directory

```
fs::dir_create(here::here("plots", "pulitzers"))
```

Create file paths

```
newspapers <- str_replace_all(tolower(final_plots$newspaper), " ", "_")  
paths <- here::here("plots", "pulitzers", glue("{newspapers}.png"))
```

Challenge

- Use a `map()` family function to loop through `paths` and `final_plots$plots` to save all plots.
- Render (knit) your file. What do you notice?

03:00

walk()

Just like `map()`, we have parallel variants of `walk()`, including, `walk2()`, and `pwalk()`

These work just like `map()` but don't print to the screen

Try replacing your prior code with a `walk()` version.

How does the rendered output change?

02:00

Save plots

```
walk2(paths, final_plots$plots, ggsave,  
      width = 9.5,  
      height = 6.5,  
      dpi = 500)
```


Unlike `map()` and its variants which always return a fixed object type (list for `map()`, integer vector for `map_int()`, etc), the `modify()` family always returns the same type as the input object.

map VS modify

map

```
map(mtcars, ~as.numeric(scale(.x)))
```

```
## $mpg
## [1] 0.15088482 0.15088482 0.44954345 0.21725341 -0.23073453 -0.3302
## [8] 0.71501778 0.44954345 -0.14777380 -0.38006384 -0.61235388 -0.4630
## [15] -1.60788262 -1.60788262 -0.89442035 2.04238943 1.71054652 2.2912
## [22] -0.76168319 -0.81145962 -1.12671039 -0.14777380 1.19619000 0.9804
## [29] -0.71190675 -0.06481307 -0.84464392 0.21725341
##
## $cyl
## [1] -0.1049878 -0.1049878 -1.2248578 -0.1049878 1.0148821 -0.1049878
## [9] -1.2248578 -0.1049878 -0.1049878 1.0148821 1.0148821 1.0148821
## [17] 1.0148821 -1.2248578 -1.2248578 -1.2248578 -1.2248578 1.0148821
## [25] 1.0148821 -1.2248578 -1.2248578 -1.2248578 1.0148821 -0.1049878
##
## $disp
## [1] -0.57061982 -0.57061982 -0.99018209 0.22009369 1.04308123 -0.0461
## [8] -0.67793094 -0.72553512 -0.50929918 -0.50929918 0.36371309 0.3637
## [15] 1.94675381 1.84993175 1.68856165 -1.22658929 -1.25079481 -1.2879
## [22] 0.70420401 0.59124494 0.96239618 1.36582144 -1.22416874 -0.8909
## [29] 0.97046468 -0.69164740 0.56703942 -0.88529152
##
```


modify

```
modify(mtcars, ~as.numeric(scale(.x)))
```

| ## | mpg | cyl | disp | hp | |
|------------------------|-------------|------------|-------------|-------------|-------|
| ## Mazda RX4 | 0.15088482 | -0.1049878 | -0.57061982 | -0.53509284 | 0.56 |
| ## Mazda RX4 Wag | 0.15088482 | -0.1049878 | -0.57061982 | -0.53509284 | 0.56 |
| ## Datsun 710 | 0.44954345 | -1.2248578 | -0.99018209 | -0.78304046 | 0.47 |
| ## Hornet 4 Drive | 0.21725341 | -0.1049878 | 0.22009369 | -0.53509284 | -0.96 |
| ## Hornet Sportabout | -0.23073453 | 1.0148821 | 1.04308123 | 0.41294217 | -0.83 |
| ## Valiant | -0.33028740 | -0.1049878 | -0.04616698 | -0.60801861 | -1.56 |
| ## Duster 360 | -0.96078893 | 1.0148821 | 1.04308123 | 1.43390296 | -0.72 |
| ## Merc 240D | 0.71501778 | -1.2248578 | -0.67793094 | -1.23518023 | 0.17 |
| ## Merc 230 | 0.44954345 | -1.2248578 | -0.72553512 | -0.75387015 | 0.60 |
| ## Merc 280 | -0.14777380 | -0.1049878 | -0.50929918 | -0.34548584 | 0.60 |
| ## Merc 280C | -0.38006384 | -0.1049878 | -0.50929918 | -0.34548584 | 0.60 |
| ## Merc 450SE | -0.61235388 | 1.0148821 | 0.36371309 | 0.48586794 | -0.98 |
| ## Merc 450SL | -0.46302456 | 1.0148821 | 0.36371309 | 0.48586794 | -0.98 |
| ## Merc 450SLC | -0.81145962 | 1.0148821 | 0.36371309 | 0.48586794 | -0.98 |
| ## Cadillac Fleetwood | -1.60788262 | 1.0148821 | 1.94675381 | 0.85049680 | -1.24 |
| ## Lincoln Continental | -1.60788262 | 1.0148821 | 1.84993175 | 0.99634834 | -1.11 |
| ## Chrysler Imperial | -0.89442035 | 1.0148821 | 1.68856165 | 1.21512565 | -0.68 |
| ## Fiat 128 | 2.04238943 | -1.2248578 | -1.22658929 | -1.17683962 | 0.90 |
| ## Honda Civic | 1.71054652 | -1.2248578 | -1.25079481 | -1.38103178 | 2.49 |
| ## Toyota Corolla | 2.29127162 | -1.2248578 | -1.28790993 | -1.19142477 | 1.16 |
| ## Toyota Corona | 0.23384555 | -1.2248578 | -0.89255318 | -0.72469984 | 0.19 |
| ## Dodge Challenger | -0.76168319 | 1.0148821 | 0.70420401 | 0.04831332 | -1.56 |
| ## AMC Javelin | -0.81145962 | 1.0148821 | 0.59124494 | 0.04831332 | -0.83 |
| ## Camaro Z28 | -1.12671039 | 1.0148821 | 0.96239618 | 1.43390296 | 0.24 |

```
map2(LETTERS[1:3], letters[1:3], paste0)
```

```
## [[1]]  
## [1] "Aa"  
##  
## [[2]]  
## [1] "Bb"  
##  
## [[3]]  
## [1] "Cc"
```

```
modify2(LETTERS[1:3], letters[1:3], paste0)
```

```
## [1] "Aa" "Bb" "Cc"
```


Iterating when errors are possible

Sometimes a loop will work for most cases, but return an error on a few

Often, you want to return the output you can

Alternatively, you might want to diagnose *where* the error is occurring

`purrr::safely`

Example

```
by_cyl <- mpg %>%  
  group_by(cyl) %>%  
  nest()  
by_cyl
```

```
## # A tibble: 4 x 2  
## # Groups:   cyl [4]  
##   cyl data  
##   <int> <list>  
## 1     4 <tibble[,10] [81 x 10]>  
## 2     6 <tibble[,10] [79 x 10]>  
## 3     8 <tibble[,10] [70 x 10]>  
## 4     5 <tibble[,10] [4 x 10]>
```

```
by_cyl %>%  
  mutate(mod = map(data, ~lm(hwy ~ displ + drv, data = .x)))
```

```
## Error: Problem with `mutate()` input `mod`.  
## x contrasts can be applied only to factors with 2 or more levels  
## i Input `mod` is `map(data, ~lm(hwy ~ displ + drv, data = .x))`.  
## i The error occurred in group 2: cyl = 5.
```

Safe return

- First, define safe function – note that this will work for any function

```
safe_lm <- safely(lm)
```

- Next, loop the safe function, instead of the standard function

```
safe_models <- by_cyl %>%  
  mutate(safe_mod = map(data, ~safe_lm(hwy ~ displ + drv, data =  
    safe_models
```

```
## # A tibble: 4 x 3  
## # Groups:   cyl [4]  
##   cyl data                safe_mod  
##   <int> <list>                <list>  
## 1     4 <tibble[,10] [81 x 10]> <named list [2]>  
## 2     6 <tibble[,10] [79 x 10]> <named list [2]>  
## 3     8 <tibble[,10] [70 x 10]> <named list [2]>  
## 4     5 <tibble[,10] [4 x 10]>  <named list [2]>
```

What's returned?

```
safe_models$safe_mod[[1]]
```

```
## $result
##
## Call:
## .f(formula = ..1, data = ..2)
##
## Coefficients:
## (Intercept)      displ      drv
##      37.370      -5.289      3.882
##
##
## $error
## NULL
```

```
safe_models$safe_mod[[4]]
```

```
## $result
## NULL
##
## $error
## <simpleError in `contrasts<-`(`*tmp*`, value = contr.funs[1 + isOF[nn]])
```

Inspecting

I often use `safely()` to help me de-bug. Why is it failing *there*.

First – create a new variable to filter for results with errors

```
safe_models %>%  
  mutate(error = map_lgl(safe_mod, ~!is.null(.x$error)))
```

```
## # A tibble: 4 x 4  
## # Groups:   cyl [4]  
##   cyl data          safe_mod          error  
##   <int> <list>          <list>          <lgl>  
## 1     4 <tibble[,10] [81 x 10]> <named list [2]> FALSE  
## 2     6 <tibble[,10] [79 x 10]> <named list [2]> FALSE  
## 3     8 <tibble[,10] [70 x 10]> <named list [2]> FALSE  
## 4     5 <tibble[,10] [4 x 10]> <named list [2]> TRUE
```


Inspecting the data

```
safe_models %>%  
  mutate(error = map_lgl(safe_mod, ~!is.null(.x$error))) %>%  
  filter(error) %>%  
  select(cyl, data) %>%  
  unnest(data)
```

```
## # A tibble: 4 x 11  
## # Groups:   cyl [1]  
##      cyl manufacturer model      displ  year trans      drv      cty      hwy  
##   <int> <chr>          <chr>    <dbl> <int> <chr>    <chr> <int> <int>  
## 1     5 volkswagen  jetta      2.5  2008 auto(s6)  f      21     29  
## 2     5 volkswagen  jetta      2.5  2008 manual(m5) f      21     29  
## 3     5 volkswagen  new beetle  2.5  2008 manual(m5) f      20     28  
## 4     5 volkswagen  new beetle  2.5  2008 auto(s6)  f      20     29
```

The **displ** and **drv** variables are constant, so no relation can be estimated.

Pull results that worked

```
safe_models %>%  
  mutate(results = map(safe_mod, "result"))
```

```
## # A tibble: 4 x 4  
## # Groups:   cyl [4]  
##   cyl data          safe_mod      results  
##   <int> <list>          <list>      <list>  
## 1     4 <tibble[,10] [81 x 10]> <named list [2]> <lm>  
## 2     6 <tibble[,10] [79 x 10]> <named list [2]> <lm>  
## 3     8 <tibble[,10] [70 x 10]> <named list [2]> <lm>  
## 4     5 <tibble[,10] [4 x 10]> <named list [2]> <NULL>
```

Now we can `broom::tidy()` or whatever

Notice that there is no `cyl == 5`.

```
safe_models %>%  
  mutate(results = map(safe_mod, "result"),  
         tidied = map(results, broom::tidy)) %>%  
  select(cyl, tidied) %>%  
  unnest(tidied)
```

```
## # A tibble: 11 x 6  
## # Groups:   cyl [3]  
##   cyl term          estimate std.error statistic    p.value  
##   <int> <chr>          <dbl>    <dbl>    <dbl>    <dbl>  
## 1     4 (Intercept)  37.37023  3.537572  10.56381 1.052943e-16  
## 2     4 displ      -5.288562  1.436068  -3.682668 4.235795e- 4  
## 3     4 drvf        3.882134  0.9971876  3.893083 2.073699e- 4  
## 4     6 (Intercept)  27.96536  2.347630  11.91217 5.718039e-19  
## 5     6 displ      -2.333261  0.6373304 -3.660991 4.651570e- 4  
## 6     6 drvf        4.570840  0.6012367  7.602397 6.789988e-11  
## 7     6 drvr        6.384355  1.229277  5.193585 1.713129e- 6  
## 8     8 (Intercept)  14.82265  2.887289  5.133759 2.708515e- 6  
## 9     8 displ        0.3060487 0.5719058  0.5351383 5.943528e- 1  
## 10    8 drvf        8.555294  2.679129  3.193311 2.156229e- 3  
## 11    8 drvr        3.709336  0.7319048  5.068058 3.473594e- 6
```

When else might we use this?

Any sort of web scraping – pages change and URLs don't always work

Example

```
library(rvest)
links <- list(
  "https://en.wikipedia.org/wiki/FC_Barcelona",
  "https://nosuchpage",
  "https://en.wikipedia.org/wiki/Rome"
)
pages <- map(links, ~{
  Sys.sleep(0.1)
  read_html(.x)
})
```

```
## Error in open.connection(x, "rb"): Failed to connect to nosuchpage port
```

The problem

I can't connect to <https://nosuchpage> because it doesn't exist

BUT

That also means I can't get *any* of my links because *one* page errored (imagine it was 1 in 1,000 instead of 1 in 3)

safely() to the rescue

Safe version

```
safe_read_html <- safely(read_html)
pages <- map(links, ~{
  Sys.sleep(0.1)
  safe_read_html(.x)
})
str(pages)
```

```
## List of 3
## $ :List of 2
## ..$ result:List of 2
## .. ..$ node:<externalptr>
## .. ..$ doc :<externalptr>
## .. ..- attr(*, "class")= chr [1:2] "xml_document" "xml_node"
## ..$ error : NULL
## $ :List of 2
## ..$ result: NULL
## ..$ error :List of 2
## .. ..$ message: chr "Failed to connect to nosuchpage port 443: Operati
## .. ..$ call : language open.connection(x, "rb")
## .. ..- attr(*, "class")= chr [1:3] "simpleError" "error" "condition"
## $ :List of 2
## ..$ result:List of 2
## .. ..$ node:<externalptr>
## .. ..$ doc :<externalptr>
## .. ..- attr(*, "class")= chr [1:2] "xml_document" "xml_node"
```

Non-results

In a real example, we'd probably want to double-check the pages where we got no results

```
errors <- map_lgl(pages, ~!is.null(.x$error))
links[errors]
```

```
## [[1]]
## [1] "https://nosuchpage"
```


reduce

Reducing a list

The `map()` family of functions will always return a vector the same length as the input

`reduce()` will collapse or reduce the list to a single element

Example

```
l <- list(  
  c(1, 3),  
  c(1, 5, 7, 9),  
  3,  
  c(4, 8, 12, 2)  
)
```

```
reduce(l, sum)
```

```
## [1] 55
```

What's going on?

The code `reduce(l, sum)` is the same as

```
sum(l[[4]], sum(l[[3]], sum(l[[1]], l[[2]])))
```

```
## [1] 55
```

Or slidghlty differently

```
first_sum <- sum(l[[1]], l[[2]])  
second_sum <- sum(first_sum, l[[3]])  
final_sum <- sum(second_sum, l[[4]])  
final_sum
```

```
## [1] 55
```

Why might you use this?

What if you had a list of data frames like this

```
l_df <- list(  
  tibble(id = 1:3, score = rnorm(3)),  
  tibble(id = 1:5, treatment = rbinom(5, 1, .5)),  
  tibble(id = c(1, 3, 5, 7), other_thing = rnorm(4))  
)
```

We can join these all together with a single loop – we want the output to be of length 1!

```
reduce(l_df, full_join)
```

```
## # A tibble: 6 x 4
##   id      score treatment other_thing
##   <dbl>    <dbl>    <int>    <dbl>
## 1     1 -1.724931         0    -1.450565
## 2     2 -0.06137023        1     NA
## 3     3 -2.459853         1    -0.9331115
## 4     4  NA              0     NA
## 5     5  NA              0    -0.2372298
## 6     7  NA             NA    -0.3091048
```

Note – you have to be careful on directionality

```
reduce(l_df, left_join)
```

```
## # A tibble: 3 x 4
##   id      score treatment
##   <dbl>    <dbl>    <int>
## 1     1 -1.724931         0
## 2     2 -0.06137023        1
## 3     3 -2.459853         1
```

```
reduce(l_df, right_join)
```

```
## # A tibble: 4 x 4
##   other_thing score treatment other_thing
##   <dbl>    <dbl>    <int>    <dbl>
## 1 -1.450565 -1.724931         0    -1.450565
## 2  NA        3 -2.459853         1    -0.9331115
## 3 -0.9331115  NA              0    -0.2372298
## 4     7  NA             NA    -0.3091048
```

More common

You probably just want to `bind_rows()`

```
l_df2 <- list(  
  tibble(id = 1:3, scid = 1, score = rnorm(3)),  
  tibble(id = 1:5, scid = 2, score = rnorm(5)),  
  tibble(id = c(1, 3, 5, 7), scid = 3, score = rnorm(4))  
)  
reduce(l_df2, bind_rows)
```

```
## # A tibble: 12 x 3  
##       id   scid   score  
##   <dbl> <dbl>   <dbl>  
## 1     1     1  1.457574  
## 2     2     1  1.468945  
## 3     3     1  0.3358508  
## 4     1     2  0.4203801  
## 5     2     2  1.142734  
## 6     3     2  0.4678470  
## 7     4     2  0.05336338  
## 8     5     2 -0.01651775  
## 9     1     3  0.02608221  
## 10    3     3  2.430486  
## 11    5     3  0.2613433  
## 12    7     3  1.026262
```

Wrap up

- Lots more to `{purrr}` but we've covered a lot
- Functional programming can *really* help your efficiency, and even if it slows you down initially, I'd recommend always striving toward it, because it will ultimately be a huge help.

Questions?

If we have any time left – let's work on the homework

Next time

Functions

Beginning next class, the focus of the course will shift