

# Functions: Part 1

---

Daniel Anderson

Week 6, Class 1

# Agenda

---

- Everything is a function
- Components of a function
- Function workflows

# Learning objectives

---

- Understand and be able to fluently refer to the three fundamental components of a function
- Understand the workflows that often lead to writing functions, and how you iterate from interactive work to writing a function
- Be able to write a few basic functions

# Functions

---

Anything that carries out an operation in R is a function. For example

```
## [1] 8
```

The `+` is a function (what's referred to as an *infix* function).

Any ideas on how we could re-write the above to make it look more "function"-y?

```
## [1] 8
```

# What about this?

---

```
## [1] 15
```

```
## [1] 15
```

or

```
## [1] 15
```

# What's going on here?

---

- The `+` operator is a function that takes two arguments (both numeric), which it sums.
- The following are also the same (minus what's being assigned)

```
## [1] 7
```

```
## [1] 5
```

Everything is a function!

# Being devious

---

- Want to introduce a devious bug? Redefine `+`

```
##  
## FALSE TRUE  
##      3   497
```

```
## Warning in rm(`+`, envir = globalenv()): object '+' not found
```

```
##  
## FALSE TRUE  
##      6   494
```

# Tricky...

---

Functions are also (usually) objects!

```
##  
## Call:  
## a(formula = hp ~ drat + wt, data = mtcars)  
##  
## Coefficients:  
## (Intercept)          drat           wt  
##      -27.782         5.354        48.244
```



# What does this all mean?

---

- Anything that carries out ANY operation in R is a function
- Functions are generally, but not always, stored in an object (otherwise known as binding the function to a name)

# Anonymous functions

---

- The function for computing the mean is bound the name `mean`
- When running things through loops, you may often want to apply a function without binding it to a name

## Example

```
##   mpg   cyl  disp    hp  drat    wt  qsec    vs  am  gear  carb
##   25     3   272    22   22    29   30     2   2    3     6
```

# Another possibility

---

- If you have a bunch of functions, you might consider storing them all in a list.
- You can then access the functions in the same way you would subset any list

*This is kind of weird...*

```
## [1] 25
```

```
## [1] 50
```

```
## [1] 400
```

# What does this imply?

---

- If we can store functions in a vector (list), then we can loop through the vector just like any other!

```
##           n      n_miss   n_valid      mean      sd
## 32.000000  0.000000 32.000000 20.090625  6.026948
```

Careful though

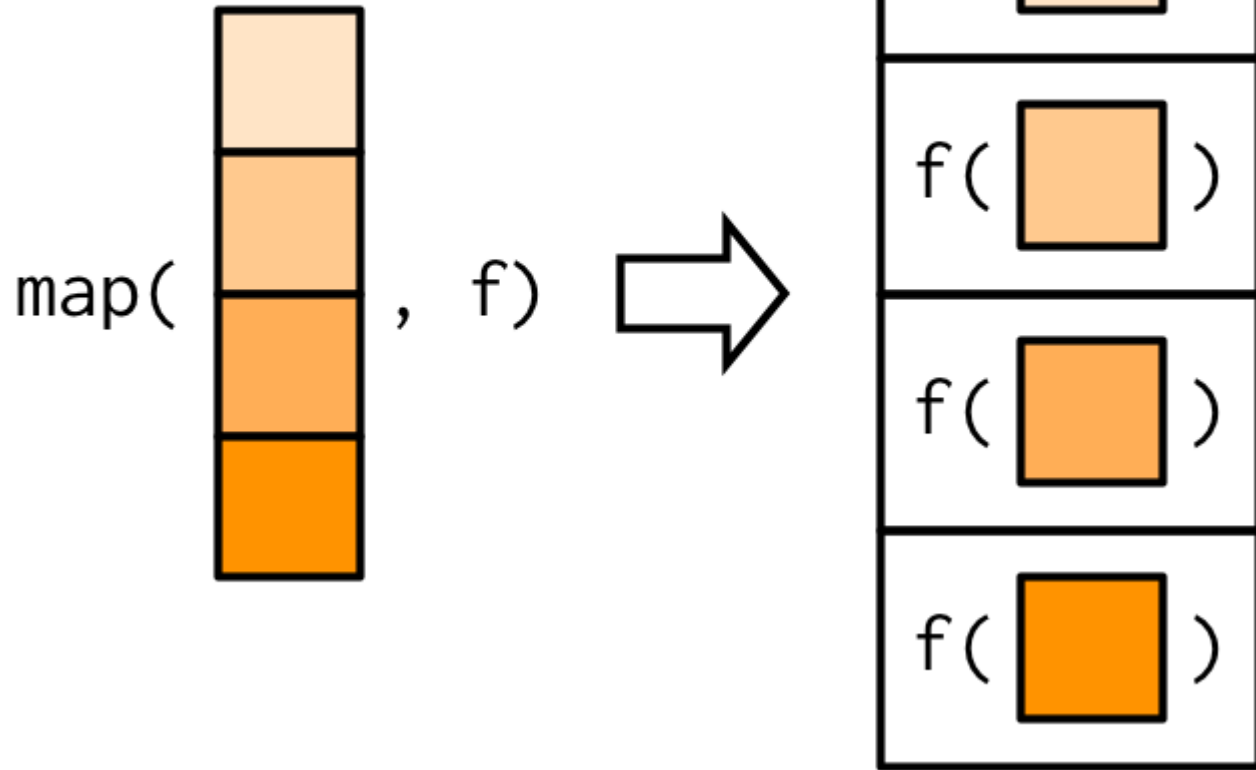
This doesn't work

```
## Error: Can't pluck from a builtin
```

Why?

Remember what {purrr}  
does

---



# With `map_df`

---

What if we wanted this for all columns?



# Challenge

---

- Can you extend the previous looping to supply the summary for every column? Hint: You'll need to make a nested loop (loop each function through each column)

```
## # A tibble: 11 x 6
##   column      n n_miss n_valid      mean      sd
##   <chr> <int> <int> <int>    <dbl>    <dbl>
## 1 mpg      32      0      32  20.09062  6.026948
## 2 cyl      32      0      32   6.1875   1.785922
## 3 disp     32      0      32 230.7219 123.9387
## 4 hp       32      0      32 146.6875  68.56287
## 5 drat     32      0      32   3.596562 0.5346787
## 6 wt       32      0      32   3.21725   0.9784574
## 7 qsec     32      0      32  17.84875   1.786943
## 8 vs       32      0      32   0.4375    0.5040161
## 9 am       32      0      32   0.40625   0.4989909
## 10 gear    32      0      32   3.6875    0.7378041
## 11 carb    32      0      32   2.8125    1.615200
```

05:00

# Maybe easier

---

- Often when you get into nested loops, it's easier to turn (at least) one of the loops into a function.

Now we can just loop this function through each column



# Wrap the whole thing in a function

---

```
## # A tibble: 6 x 6
##   column      n n_miss n_valid      mean      sd
##   <chr>   <int>  <int>   <int>   <dbl>   <dbl>
## 1 Ozone    153    37    116 NA      NA
## 2 Solar.R  153     7    146 NA      NA
## 3 Wind     153     0    153  9.957516  3.523001
## 4 Temp     153     0    153  77.88235  9.465270
## 5 Month    153     0    153  6.993464  1.416522
## 6 Day      153     0    153  15.80392  8.864520
```

Notice the missing data. Why? What should we do?

# Function

# components


---

# Three components

---

- `body()`
- `formals()`
- `environment()` (we won't focus so much here for now)

```
poly <- function(x, power) {  
  z <- x^power  
  return(z)  
}
```



Body

Formals

# Formals

---

- The arguments supplied to the function
- What's one way to identify the formals for a function – say, `lm`?

`?`: Help documentation!

Alternative – use a function!

```
## $formula
##
##
## $data
##
##
## $subset
##
##
## $weights
##
##
## $na.action
```

# How do you see the body?

---

- In RStudio: Super (command on mac, cntrl on windows) + click!

[demo]

- Alternative – just print to screen



# Or use **body**

---

```
## {
##   ret.x <- x
##   ret.y <- y
##   cl <- match.call()
##   mf <- match.call(expand.dots = FALSE)
##   m <- match(c("formula", "data", "subset", "weights", "na.action",
##               "offset"), names(mf), 0L)
##   mf <- mf[c(1L, m)]
##   mf$drop.unused.levels <- TRUE
##   mf[[1L]] <- quote(stats::model.frame)
##   mf <- eval(mf, parent.frame())
##   if (method == "model.frame")
##     return(mf)
##   else if (method != "qr")
##     warning(gettextf("method = '%s' is not supported. Using 'qr'",
##                     method), domain = NA)
##   mt <- attr(mf, "terms")
##   y <- model.response(mf, "numeric")
##   w <- as.vector(model.weights(mf))
##   if (!is.null(w) && !is.numeric(w))
##     stop("'weights' must be a numeric vector")
##   offset <- model.offset(mf)
##   mlm <- is.matrix(y)
##   ny <- if (mlm)
##     nrow(y)
##   else length(y)
##   if (!is.null(offset)) {
```

# Environment

---

- As I mentioned, we won't focus on this too much, but if you get deep into programming it's pretty important

```
## <environment: 0x7fd29d8ca2d0>
```

```
## <environment: namespace:stats>
```

# Why this matters

---

What will the following return?

```
## [1] 20
```

# What will this return?

---

```
## [1] 3
```

# Last one

---

What do each of the following return?

```
## [1] 3
```

```
## Error in eval(expr, envir, enclos): object 'y' not found
```

# Environment summary

---

- The previous examples are part of *lexical scoping*.
- Generally, you won't have to worry too much about it
- If you end up with unexpected results, this could be part of why

# Scoping

---

- Part of what's interesting about these scoping rules is that your functions can, and very often do, depend upon things in your global workspace, or your specific environment.
- If this is the case, the function will be a "one-off", and unlikely to be useful in any other script

# Example 1

---

## Extracting information

- This is a real example



# Example 2

---

Reading in data

# Simple example

---

Pull out specific coefficients

```
## # A tibble: 3 x 3
## # Groups:   cyl [3]
##     cyl data                                model
##   <dbl> <list>                                <list>
## 1     6 <tibble[,10] [7 x 10]> <lm>
## 2     4 <tibble[,10] [11 x 10]> <lm>
## 3     8 <tibble[,10] [14 x 10]> <lm>
```

# Pull a specific coef

---

Find the solution for one model

```
## (Intercept)          disp             hp          drat
## 6.284507434 0.026354099 0.006229086 2.193576546
```

```
##          disp
## 0.0263541
```

```
## (Intercept)
## 6.284507
```

# Generalize it

---

```
## # A tibble: 3 x 7
## # Groups:   cyl [3]
##   cyl data          model intercept      disp      h
##   <dbl> <list>          <list>      <dbl>      <dbl>      <dbl>
## 1     6 <tibble[,10] [7 x 10]> <lm>      6.284507  0.02635410  0.00622908
## 2     4 <tibble[,10] [11 x 10]> <lm>     46.08662 -0.1225361 -0.04937771
## 3     8 <tibble[,10] [14 x 10]> <lm>     19.00162 -0.01671461 -0.02140236
```

# Make it more flexible

---

- Since the intercept is a little difficult to pull out, we could have it return that by default.

```
## # A tibble: 3 x 4
## # Groups:   cyl [3]
##   cyl data          model intercept
##   <dbl> <list>          <list>      <dbl>
## 1     6 <tibble[,10] [7 x 10]> <lm>      6.284507
## 2     4 <tibble[,10] [11 x 10]> <lm>     46.08662
## 3     8 <tibble[,10] [14 x 10]> <lm>     19.00162
```

# Return all coefficients

---

```
## # A tibble: 3 x 4
## # Groups:   cyl [3]
##       cyl data          model  coefs
##   <dbl> <list>          <list> <list>
## 1     6 <tibble[,10] [7 x 10]> <lm>    <df[,2] [4 x 2]>
## 2     4 <tibble[,10] [11 x 10]> <lm>    <df[,2] [4 x 2]>
## 3     8 <tibble[,10] [14 x 10]> <lm>    <df[,2] [4 x 2]>
```

```
## # A tibble: 12 x 5
## # Groups:   cyl [3]
##       cyl data                model coefficient      estimate
##   <dbl> <list>                <list> <chr>          <dbl>
## 1     6 <tibble[,10] [7 x 10]> <lm> (Intercept)    6.284507
## 2     6 <tibble[,10] [7 x 10]> <lm> disp         0.02635410
## 3     6 <tibble[,10] [7 x 10]> <lm> hp           0.006229086
## 4     6 <tibble[,10] [7 x 10]> <lm> drat         2.193577
## 5     4 <tibble[,10] [11 x 10]> <lm> (Intercept) 46.08662
## 6     4 <tibble[,10] [11 x 10]> <lm> disp        -0.1225361
## 7     4 <tibble[,10] [11 x 10]> <lm> hp          -0.04937771
## 8     4 <tibble[,10] [11 x 10]> <lm> drat        -0.6041857
## 9     8 <tibble[,10] [14 x 10]> <lm> (Intercept) 19.00162
## 10    8 <tibble[,10] [14 x 10]> <lm> disp        -0.01671461
## 11    8 <tibble[,10] [14 x 10]> <lm> hp          -0.02140236
## 12    8 <tibble[,10] [14 x 10]> <lm> drat         2.006011
```

# Slightly nicer

---

```
## # A tibble: 12 x 3
## # Groups:   cyl [3]
##       cyl coefficient      estimate
##   <dbl> <chr>          <dbl>
## 1     6 (Intercept)  6.284507
## 2     6 disp         0.02635410
## 3     6 hp          0.006229086
## 4     6 drat        2.193577
## 5     4 (Intercept) 46.08662
## 6     4 disp        -0.1225361
## 7     4 hp          -0.04937771
## 8     4 drat        -0.6041857
## 9     8 (Intercept) 19.00162
## 10    8 disp        -0.01671461
## 11    8 hp          -0.02140236
## 12    8 drat        2.006011
```



# Create nice table

---

```
## # A tibble: 3 x 5
## # Groups:   cyl [3]
##   cyl `(Intercept)`      disp          hp      drat
##   <dbl>         <dbl>      <dbl>      <dbl>      <dbl>
## 1     4      46.08662 -0.1225361 -0.04937771 -0.6041857
## 2     6       6.284507  0.02635410  0.006229086  2.193577
## 3     8      19.00162 -0.01671461 -0.02140236  2.006011
```

When to write  
a function?

---

# Example

---

```
## # A tibble: 10 x 4
##           a           b           c           d
##       <dbl>       <dbl>       <dbl>       <dbl>
## 1 305.6438    295.7304    54.00421    168.3175
## 2  15.29527   442.9968   -167.1963    205.7256
## 3 154.4693   -108.3291    74.21240    255.2655
## 4 194.9294    58.18168   282.2012     8.661044
## 5 160.6402    80.00180   384.2790    175.7433
## 6  84.08132   195.3926    35.42963   -157.5513
## 7 326.7283    57.36206    61.40959   -17.66885
## 8  85.80114  -298.4683   -164.4745   -27.63614
## 9 402.7636   -266.0700    169.0146   -262.1311
## 10 90.59289   298.0170     4.000769   105.4184
```

# Rescale each column to 0/1

---

Do it for one column

```
## # A tibble: 10 x 4
##           a           b           c           d
##       <dbl>       <dbl>       <dbl>       <dbl>
## 1 0.7493478 295.7304    54.00421    168.3175
## 2 0         442.9968  -167.1963    205.7256
## 3 0.3591881 -108.3291    74.21240    255.2655
## 4 0.4636099  58.18168   282.2012     8.661044
## 5 0.3751145  80.00180   384.2790    175.7433
## 6 0.1775269 195.3926    35.42963  -157.5513
## 7 0.8037639  57.36206    61.40959   -17.66885
## 8 0.1819655 -298.4683  -164.4745   -27.63614
## 9 1         -266.0700   169.0146  -262.1311
## 10 0.1943323 298.0170     4.000769   105.4184
```

# Do it for all columns

---

```
## # A tibble: 10 x 4
##           a           b           c           d
##       <dbl>       <dbl>       <dbl>       <dbl>
## 1 0.7493478 0.8013846 0.4011068 0.8319510
## 2 0          1          0          0.9042516
## 3 0.3591881 0.2564372 0.4377506 1
## 4 0.4636099 0.4810071 0.8149005 0.5233744
## 5 0.3751145 0.5104355 1          0.8463031
## 6 0.1775269 0.6660608 0.3674252 0.2021270
## 7 0.8037639 0.4799017 0.4145351 0.4724852
## 8 0.1819655 0          0.004935493 0.4532209
## 9 1          0.04369494 0.6096572 0
## 10 0.1943323 0.8044685 0.3104346 0.7103825
```

# An alternative

---

- What's an alternative we could use *without* writing a function?

```
## # A tibble: 10 x 4
##           a           b           c           d
##       <dbl>       <dbl>       <dbl>       <dbl>
## 1 0.7493478 0.8013846 0.4011068 0.8319510
## 2 0          1          0          0.9042516
## 3 0.3591881 0.2564372 0.4377506 1
## 4 0.4636099 0.4810071 0.8149005 0.5233744
## 5 0.3751145 0.5104355 1          0.8463031
## 6 0.1775269 0.6660608 0.3674252 0.2021270
## 7 0.8037639 0.4799017 0.4145351 0.4724852
## 8 0.1819655 0          0.004935493 0.4532209
## 9 1          0.04369494 0.6096572 0
## 10 0.1943323 0.8044685 0.3104346 0.7103825
```

# Another alternative

---

Write a function

- What are the arguments going to be?
- What will the body be?

Arguments

- One formal argument – A numeric vector to rescale

# Body

---

- You try first

02:00



# Create the function

---

Test it!

```
## [1] 0.0 0.5 1.0
```

```
## [1] 0.0 0.1 0.2 0.3 0.4 0.5 0.6 0.7 0.8 0.9 1.0
```

# Make it cleaner

---

- There's nothing inherently "wrong" about the prior function, but it is a bit hard to read
- How could we make it easier to read?
  - Remove missing data once (rather than every time)
  - Don't calculate things multiple times

# A little cleaned up

---

Test it!

```
## [1] 0.0 0.5 1.0
```

```
## [1] 0.0 0.1 0.2 0.3 0.4 0.5 0.6 0.7 0.8 0.9 1.0
```

Make sure they give the same  
output

```
## [1] TRUE
```

```
## [1] TRUE
```

# Final solution

---

Could use **modify** here too

```
## # A tibble: 10 x 4
##       a          b          c          d
##   <dbl>    <dbl>    <dbl>    <dbl>
## 1 0.7493478 0.8013846 0.4011068 0.8319510
## 2 0          1          0          0.9042516
## 3 0.3591881 0.2564372 0.4377506 1
## 4 0.4636099 0.4810071 0.8149005 0.5233744
## 5 0.3751145 0.5104355 1          0.8463031
## 6 0.1775269 0.6660608 0.3674252 0.2021270
## 7 0.8037639 0.4799017 0.4145351 0.4724852
## 8 0.1819655 0          0.004935493 0.4532209
## 9 1          0.04369494 0.6096572 0
## 10 0.1943323 0.8044685 0.3104346 0.7103825
```

# Getting more complex

---

- What if you want a function to behave differently depending on the input?

Add conditions

# Lots of conditions?

---

# Easy example

---

- Given a vector, return the mean if it's numeric, and **NULL** otherwise



# Test it

---

```
## [1] 0.1855869
```

```
## NULL
```

# Mean for all numeric columns

---

- The prior function can now be used within a new function to calculate the mean of all columns of a data frame that are numeric

##	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
## 1	5.1	3.5	1.4	0.2	setosa
## 2	4.9	3.0	1.4	0.2	setosa
## 3	4.7	3.2	1.3	0.2	setosa
## 4	4.6	3.1	1.5	0.2	setosa
## 5	5.0	3.6	1.4	0.2	setosa
## 6	5.4	3.9	1.7	0.4	setosa

##	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width
## 1	5.843333	3.057333	3.758	1.199333

# We have a problem though!

---

```
##      Ozone  Solar.R  Wind  Temp  Month  Day
## 1      41      190   7.4   67     5    1
## 2      36      118   8.0   72     5    2
## 3      12      149  12.6   74     5    3
## 4      18      313  11.5   62     5    4
## 5      NA       NA  14.3   56     5    5
## 6      28       NA  14.9   66     5    6
```

```
##      Ozone  Solar.R      Wind      Temp      Month      Day
## 1      NA       NA  9.957516  77.88235  6.993464  15.80392
```

Why is this happening?

How can we fix it?

# Easiest way in this case . . .

---

Pass the dots!

Redefine `means2`

# Reefine means\_df

---

##	Ozone	Solar.R	Wind	Temp	Month	Day
## 1	NA	NA	9.957516	77.88235	6.993464	15.80392

##	Ozone	Solar.R	Wind	Temp	Month	Day
## 1	42.12931	185.9315	9.957516	77.88235	6.993464	15.80392

# Next time

---

Functions: Part 2