

# Developing Your First R Package

---

Daniel Anderson

Week 10

# Agenda

---

# Want to follow along?

---

If you'd like to follow along, please make sure you have the following packages installed

```
install.packages(c("tidyverse", "devtools", "esvis",  
                  "roxygen2", "usethis"))
```

# Bundle your functions

---

Once you've written more than one function, you may want to bundle them. There are two general ways to do this:

source?

Write a package



# Why avoid **sourceing**

---

- Documentation is generally more sparse
- Directory issues
  - Which leads to reproducibility issues
  - This is also less of an issue if you're using RStudio Projects and {here}

# More importantly

---

Bundling functions into a package is not that hard!



# My journey with {esvis}

---

My first CRAN package

# Background

---

## Effect sizes

### Standardized mean differences

- Assumes reasonably normally distributed distributions (mean is a good indicator of central tendency)
- Differences in means may not reflect differences at all points in scale if variances are different
- Substantive interest may also lie with differences at other points in the distribution.



# Varying differences

---

## Quick simulated example

```
library(tidyverse)
common_var <- tibble(low = rnorm(1000, 10, 1),
                     high = rnorm(1000, 12, 1),
                     var = "common")
diff_var <- tibble(low = rnorm(1000, 10, 1),
                   high = rnorm(1000, 12, 2),
                   var = "diff")
d <- bind_rows(common_var, diff_var)
head(d)
```

```
## # A tibble: 6 x 3
##       low      high var
##   <dbl>   <dbl> <chr>
## 1  9.549110 13.50911 common
## 2  9.350832 12.93887 common
## 3 10.40181  12.20612 common
## 4 10.81037  10.52069 common
## 5 10.79728  12.71338 common
## 6  9.672677 10.68060 common
```

# Restructure for plotting


---

```
d <- d %>%  
  pivot_longer(  
    -var,  
    names_to = "group",  
    values_to = "value"  
  )  
d
```

```
## # A tibble: 4,000 x 3  
##   var      group      value  
##   <chr>   <chr>      <dbl>  
## 1 common low      9.549110  
## 2 common high    13.50911  
## 3 common low      9.350832  
## 4 common high    12.93887  
## 5 common low     10.40181  
## 6 common high    12.20612  
## 7 common low     10.81037  
## 8 common high    10.52069  
## 9 common low     10.79728  
## 10 common high   12.71338  
## # ... with 3,990 more rows
```

# Plot the distributions

```
ggplot(d, aes(value, fill = group)) +  
  geom_density(alpha = 0.7,  
               color = "gray40") +  
  facet_wrap(~var) +  
  scale_fill_brewer(palette = "Set3")
```



# Binned effect sizes

---

1. Cut the distributions into ***n*** bins (based on percentiles)
2. Calculate the mean difference between paired bins
3. Divide each mean difference by the overall pooled standard deviation

$$d_{[i]} = \frac{\bar{X}_{foc[i]} - \bar{X}_{ref[i]}}{\sqrt{\frac{(n_{foc}-1)Var_{foc} + (n_{ref}-1)Var_{ref}}{n_{foc} + n_{ref} - 2}}}$$

visualize it!

# Back to the simulated example

---

```
common <- filter(d, var == "common")  
diff    <- filter(d, var == "diff")
```

```
library(esvis)
binned_es(common, value ~ group)
```

```
## # A tibble: 6 x 11
##       q  qtile_lb  qtile_ub group_ref group_foc mean_diff length length1
##   <dbl>    <dbl>    <dbl> <chr>    <chr>    <dbl>  <int>  <int>
## 1     1  0        0.3333333 high    low     -2.048922 1000 1000
## 2     2 0.3333333 0.6666667 high    low     -2.080733 1000 1000
## 3     3 0.6666667 1          high    low     -1.992408 1000 1000
## 4     1  0        0.3333333 low     high     2.048922 1000 1000
## 5     2 0.3333333 0.6666667 low     high     2.080733 1000 1000
## 6     3 0.6666667 1          low     high     1.992408 1000 1000
```

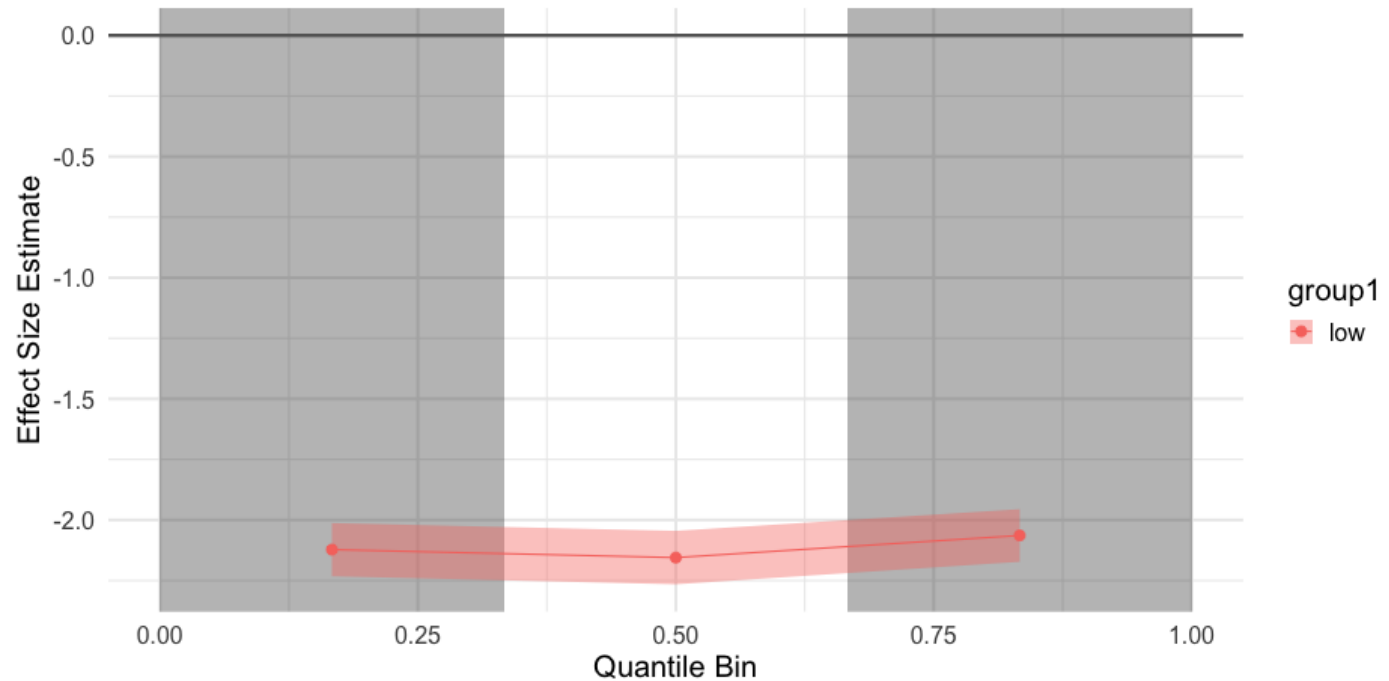
```
binned_es(diff, value ~ group)
```

```
## # A tibble: 6 x 11
##       q  qtile_lb  qtile_ub group_ref group_foc mean_diff length length1
##   <dbl>    <dbl>    <dbl> <chr>    <chr>    <dbl>  <int>  <int>
## 1     1  0        0.3333333 high    low     -1.098328 1000 1000
## 2     2 0.3333333 0.6666667 high    low     -1.967218 1000 1000
## 3     3 0.6666667 1          high    low     -3.133526 1000 1000
## 4     1  0        0.3333333 low     high     1.098328 1000 1000
## 5     2 0.3333333 0.6666667 low     high     1.967218 1000 1000
## 6     3 0.6666667 1          low     high     3.133526 1000 1000
```

# Visualize it

## Common Variance

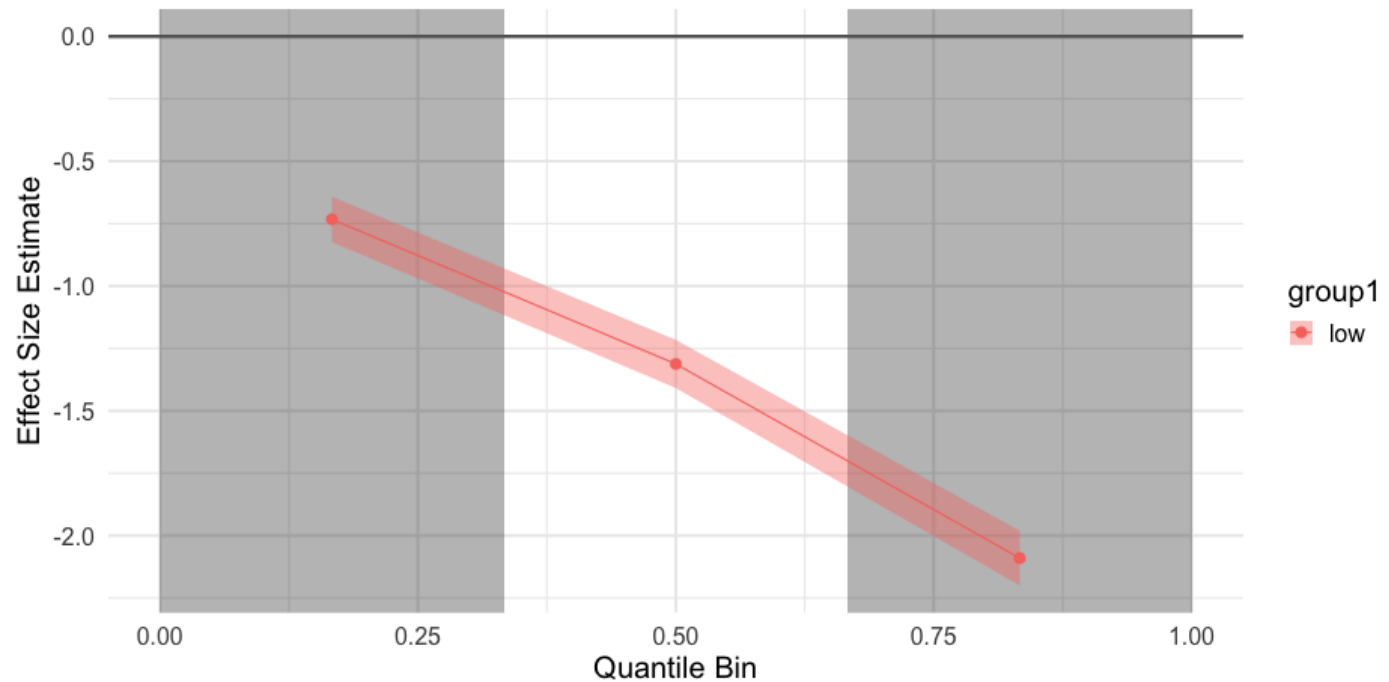
```
binned_plot(common, value ~ group)
```



# Visualize it

## Different Variance

```
binned_plot(diff, value ~ group)
```






# Wait a minute...

---

- The *esvis* package will (among other things) calculate and visually display binned effect sizes.
- But how did we get from an idea, to functions, to a package?



Taking a step  
back

---

# Package Creation

---

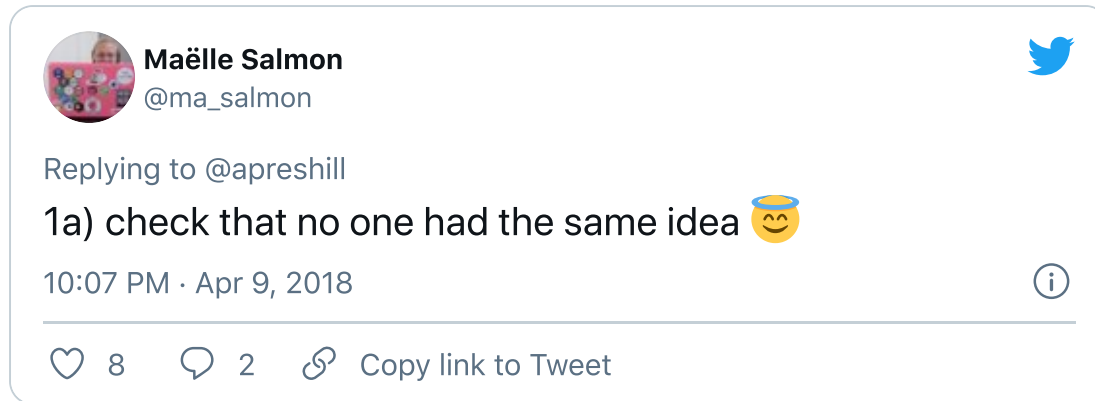
## The (or rather a) recipe

1. Come up with ~~a brilliant~~ an idea
  - can be boring and mundane but just something you do a lot
2. Write a function! or more likely, a set of functions
3. Create package skelton
4. Document your function
5. Install/fiddle/install
6. Write tests for your functions
7. Host your package somewhere public (GitHub is probably best) and promote it – leverage the power of open source!

Use tools to automate

# A really good point

---



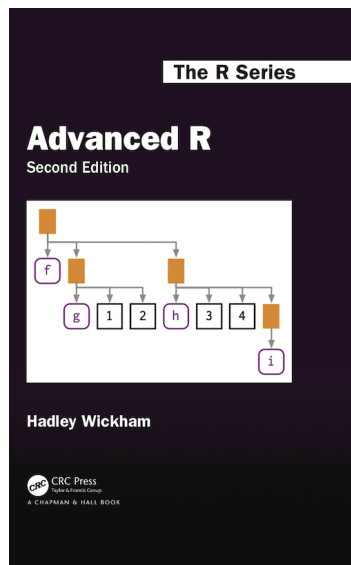
And some further recommendations/good advice

# Some resources


---

We surely won't get through everything. In my mind, the best resources are:

Advanced R



R Packages



# Our package

---

We're going to write a package today! Let's keep it really simple...

1. Idea (which we've actually used before): Report basic descriptive statistics for a vector, `x`: `N`, `n-valid`, `n-missing`, `mean`, and `sd`.

# Our function

---

- Let's have it return a data frame
- What will be the formal arguments?
- What will the body look like?

Want to give it a go?

# The approach I took...

---

```
describe <- function(data, column_name) {  
  x <- data[[column_name]]  
  
  nval <- length(na.omit(x))  
  nmiss <- sum(is.na(x))  
  mn <- mean(x, na.rm = TRUE)  
  stdev <- sd(x, na.rm = TRUE)  
  
  out <- tibble::tibble(N           = nval + nmiss,  
                        n_valid    = nval,  
                        n_missing  = nmiss,  
                        mean       = mn,  
                        sd         = stdev)  
  
  out  
}
```



# The approach I took...

---

```
describe <- function(data, column_name) {  
  x <- data[[column_name]] # Extract just the vector to summarize  
  
  nval  <- length(na.omit(x))  
  nmiss <- sum(is.na(x))  
  mn    <- mean(x, na.rm = TRUE)  
  stdev <- sd(x, na.rm = TRUE)  
  
  out <- tibble::tibble(N          = nval + nmiss,  
                        n_valid    = nval,  
                        n_missing  = nmiss,  
                        mean       = mn,  
                        sd         = stdev)  
  
  out  
}
```

# The approach I took...

---

```
describe <- function(data, column_name) {  
  x <- data[[column_name]]  
  
  nval <- length(na.omit(x)) # Count non-missing  
  nmiss <- sum(is.na(x)) # Count missing  
  mn <- mean(x, na.rm = TRUE) # Compute mean  
  stdev <- sd(x, na.rm = TRUE) # Computer SD  
  
  out <- tibble::tibble(N      = nval + nmiss,  
                        n_valid = nval,  
                        n_missing = nmiss,  
                        mean     = mn,  
                        sd       = stdev)  
  
  out  
}
```

# The approach I took...

---

```
describe <- function(data, column_name) {  
  x <- data[[column_name]]  
  
  nval  <- length(na.omit(x))  
  nmiss <- sum(is.na(x))  
  mn    <- mean(x, na.rm = TRUE)  
  stdev <- sd(x, na.rm = TRUE)  
  
  # Compile into a df  
  out <- tibble::tibble(N      = nval + nmiss,  
                        n_valid = nval,  
                        n_missing = nmiss,  
                        mean    = mn,  
                        sd      = stdev)  
  
  out  
}
```

# The approach I took...

---

```
describe <- function(data, column_name) {  
  x <- data[[column_name]]  
  
  nval      <- length(na.omit(x))  
  nmiss <- sum(is.na(x))  
  mn      <- mean(x, na.rm = TRUE)  
  stdev   <- sd(x, na.rm = TRUE)  
  
  out <- tibble::tibble(N          = nval + nmiss,  
                        n_valid    = nval,  
                        n_missing  = nmiss,  
                        mean       = mn,  
                        sd         = stdev)  
  out # Return the table  
}
```

# Informal testing

---

```
set.seed(8675309)
df1 <- tibble(x = rnorm(100))
df2 <- tibble(var_miss = c(rnorm(1000, 10, 4), rep(NA, 27)))
describe(df1, "x")
```

```
## # A tibble: 1 x 5
##       N n_valid n_missing      mean      sd
##   <int>   <int>   <int>    <dbl>    <dbl>
## 1    100     100       0 0.05230278 0.9291437
```

```
describe(df2, "var_miss")
```

```
## # A tibble: 1 x 5
##       N n_valid n_missing      mean      sd
##   <int>   <int>   <int>    <dbl>    <dbl>
## 1   1027    1000      27 9.881107 4.090208
```

# Demo

---

Package skeleton:

- `usethis::create_package()`
- `usethis::use_r()`
- Use `roxygen2` special comments for documentation
- Run `devtools::document()`
- Install and restart, play around

# roxygen2 comments

---

## Typical arguments

- **@param**: Describe the formal arguments. State argument name and the describe it.

```
| #' @param x Vector to describe
```

- **@return**: What does the function return

```
| #' @return A tibble with descriptive  
| data
```

- **@example** or more commonly **@examples**: Provide examples of the use of your function.

- `@export`: Export your function

If you don't include `@export`, your function will be *internal*, meaning others can't access it easily.



# Other docs

---

- **NAMESPACE**: Created by **{roxygen2}**. Don't edit it. If you need to, trash it and it will be reproduced.
- **DESCRIPTION**: Describes your package (more on next slide)
- **man/**: The documentation files. Created by **{roxygen2}**. Don't edit.

# DESCRIPTION

---

Metadata about the package. Default fields for our package are

```
Package: practice
Version: 0.0.0.9000
Title: What the Package Does (One Line, Title Case)
Description: What the package does (one paragraph).
Authors@R: person("First", "Last", email = "first.last@example.com", role =
License: What license is it under?
Encoding: UTF-8
LazyData: true
ByteCompile: true
RoxygenNote: 6.0.1
```

This is where the information for `citation(package = "practice")` will come from.

Some advice – edit within RStudio, or a good text editor like sublimetext or VSCode. "Fancy" quotes and things can screw this up.

# Description File Fields

---

The 'Package', 'Version', 'License', 'Description', 'Title', 'Author', and 'Maintainer' fields are mandatory, all other fields are optional.

– Writing R Extensions

Some optional fields include

- Imports and Suggests (we'll do this in a minute).
- URL
- BugReports
- License (we'll have {usethis} create this for us).
- LazyData

# DESCRIPTION for {esvis}

---

Package: esvis

Type: Package

Title: Visualization and Estimation of Effect Sizes

Version: 0.3.1

Authors@R: person("Daniel", "Anderson", email = "daniela@uoregon.edu",  
role = c("aut", "cre"))

Description: A variety of methods are provided to estimate and visualize distributional differences in terms of effect sizes. Particular emphasis is upon evaluating differences between two or more distributions across the entire scale, rather than at a single point (e.g., differences in means). For example, Probability-Probability (PP) plots display the difference between two or more distributions, matched by their empirical CDFs (see Ho and Reardon, 2012; <doi:10.3102/1076998611411918>), allowing for examinations of where on the scale distributional differences are largest or smallest. The area under the PP curve (AUC) is an effect-size metric, corresponding to the probability that a randomly selected observation from the x-axis distribution will have a higher value than a randomly selected observation from the y-axis distribution. Binned effect size plots are also available, in which the distributions are split into bins (set by the user) and separate effect sizes (Cohen's d) are produced for each bin - again providing a means to evaluate the consistency (or lack thereof) of the difference between two or more distributions at different points on the scale. Evaluation of empirical CDFs is also provided, with built-in arguments for providing annotations to help evaluate distributional differences at specific points (e.g., semi-transparent shading). All functions take a consistent argument structure. Calculation of specific effect sizes is also possible. The

# DESCRIPTION for {esvis}

## (continued)

---

Depends:

R (>= 3.1)

Imports:

sfsmisc,  
ggplot2,  
magrittr,  
dplyr,  
rlang,  
tidyr (>= 1.0.0),  
purrr,  
Hmisc,  
tibble

URL: <https://github.com/datalorax/esvis>

BugReports: <https://github.com/datalorax/esvis/issues>

License: MIT + file LICENSE

LazyData: true

RoxygenNote: 7.0.2

Suggests:

testthat,  
viridisLite

# Demo

---

- Change the author name.
  - Add a contributor just for fun.
- Add a license. We'll go for MIT license using  
`usethis::use_mit_license("First and Last Name")`
- Install and reload.

# Declare dependencies

---

- The function **depends on** the **tibble** function within the {tibble} package.
- We have to declare this dependency

# My preferred approach

---

- Declare package dependencies:  
`usethis::use_package()`
- Create a package documentation page:  
`usethis::use_package_doc()`
  - Declare all dependencies for your package there
  - Only import the functions you need – not the entire package
    - Use `#' importFrom pkg fun_name`
- Generally won't have to worry about namespacing. The likelihood of conflicts is also reduced, so long as you don't import the full package.



# Demo

---

# Write tests!

---

- What does it mean to write tests?
  - ensure your package does what you expect it to
- Why write tests?
  - If you write a new function, and it breaks an old one, that's good to know!
  - Reduces bugs, makes your package code more robust

# How

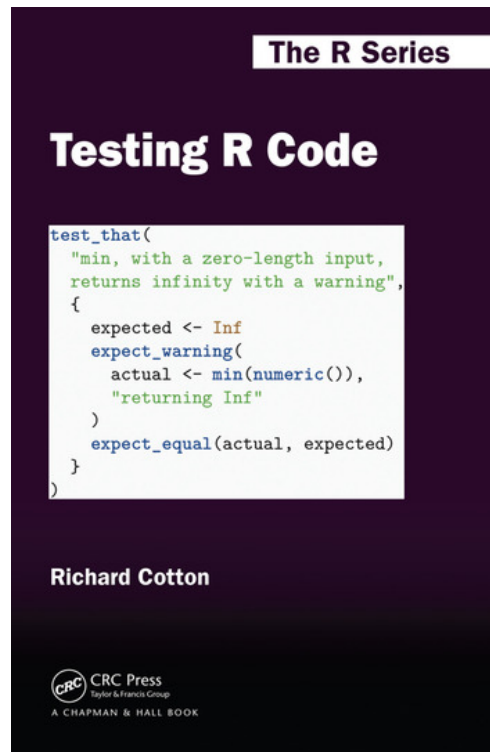
---

- `usethis::use_testthat` sets up the infrastructure
- Make assertions, e.g.: `testthat::expect_equal()`,  
`testthat::expect_warning()`,  
`testthat::expect_error()`

# Testing

---

We'll skip over testing for today, because we just don't have time to cover everything. A few good resources:



[Richie Cotton's book](#)

[r-pkgs Chapter](#)

[Karl Broman Blog Post](#)

# Check your R package

---

- Use `devtools::check()` to run the same checks CRAN will run on your R package.
  - Use `devtools::check_rhub()` to test your package on <https://builder.r-hub.io/> (several platforms and R versions)
  - Use `devtools::build_win()` to run the checks on CRAN computers.

I would not run the latter two until you're getting close to being ready to submit to CRAN.

# Patience

---

The first time, you'll likely get errors. It will probably be frustrating, but ultimately worth the effort.



Let's check  
now!

---

🎉 Hooray! 🎉

---

You have a package!





# A few other best practices

---

- Create a **README** with `usethis::use_readme_rmd`.
- Try to get your code coverage up above 80%.
- Automate wherever possible ({devtools} and {usethis} help a lot with this)
- Use the {goodpractice} package to help you package code be more robust, specifically with `goodpractice::gp()`. It will give you lots of good ideas

# A few other best practices

---

- Host on GitHub, and capitalize on integration with other systems (all free, but require registering for an account)
  - [Github Actions](#)
  - [codecov](#)

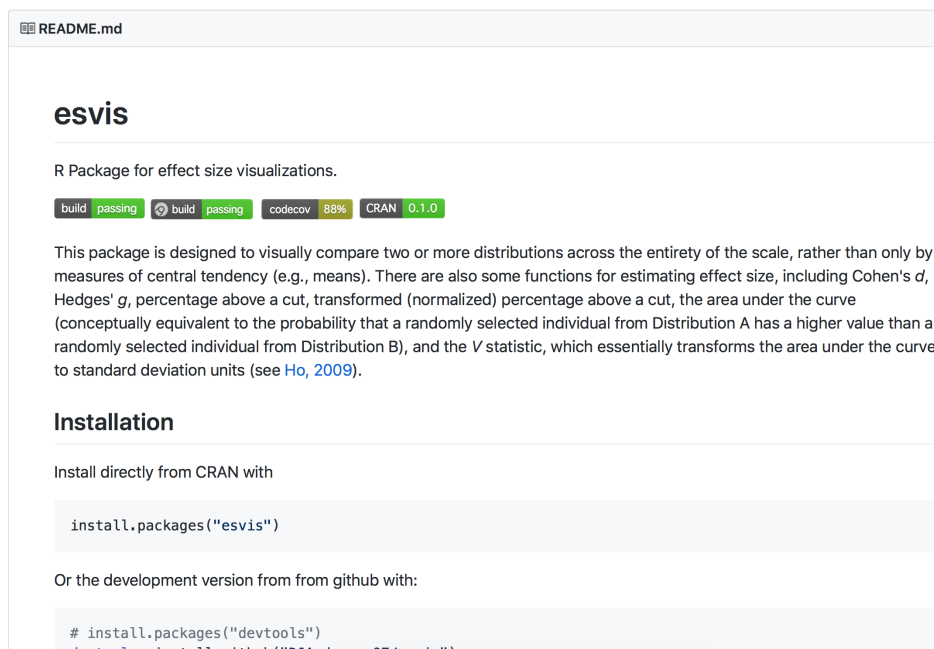
Any time left?

---

# Create a README

---

- Use standard R Markdown. Setup the infrastructure with `usethis::use_readme_rmd`.
- Write it just like a normal R Markdown doc and it should all flow into the README.



The screenshot shows a README.md file for the 'esvis' R package. The title 'esvis' is in bold. Below it, a subtitle reads 'R Package for effect size visualizations.' A row of badges follows: 'build passing', 'build passing', 'codecov 88%', and 'CRAN 0.1.0'. The main text describes the package's purpose: 'This package is designed to visually compare two or more distributions across the entirety of the scale, rather than only by measures of central tendency (e.g., means). There are also some functions for estimating effect size, including Cohen's  $d$ , Hedges'  $g$ , percentage above a cut, transformed (normalized) percentage above a cut, the area under the curve (conceptually equivalent to the probability that a randomly selected individual from Distribution A has a higher value than a randomly selected individual from Distribution B), and the  $V$  statistic, which essentially transforms the area under the curve to standard deviation units (see [Ho, 2009](#)).' A section titled 'Installation' follows. It first says 'Install directly from CRAN with' and shows the code `install.packages("esvis")`. Then it says 'Or the development version from from github with:' and shows the code `# install.packages("devtools")` followed by `devtools::install_github("D1Anderson07/esvis")`.

README.md

## esvis

R Package for effect size visualizations.

build passing build passing codecov 88% CRAN 0.1.0

This package is designed to visually compare two or more distributions across the entirety of the scale, rather than only by measures of central tendency (e.g., means). There are also some functions for estimating effect size, including Cohen's  $d$ , Hedges'  $g$ , percentage above a cut, transformed (normalized) percentage above a cut, the area under the curve (conceptually equivalent to the probability that a randomly selected individual from Distribution A has a higher value than a randomly selected individual from Distribution B), and the  $V$  statistic, which essentially transforms the area under the curve to standard deviation units (see [Ho, 2009](#)).

### Installation

Install directly from CRAN with

```
install.packages("esvis")
```

Or the development version from from github with:

```
# install.packages("devtools")
devtools::install_github("D1Anderson07/esvis")
```

# Use GitHub Actions

---

- Run `usethis::use_github_actions()` to get started.
  - Go to the Actions tab on your repo
  - Copy and paste the code to the badge into your **README**.
- Now all your code will be automatically tested each time you push! You can even test on different operating systems and different versions of R!

# codecov

---

You can test your code coverage each time you push a new commit by using `codecov`. Initialize with `usethis::use_coverage()`. Overall setup process is pretty similar to Travis CI/Appveyor.

Easily see what is/is not covered by tests!

# That's all

---

Thanks so much!